

Infra as Code

Why do it?

- Getting started with **I**nfrastructure **a**s **C**ode (IAC)
- Getting to know Terraform
- Build a sample Web Application using Terraform

What We will be covering

- Session 2 – Infrastructure as Code
 - Concepts and benefits
 - Terraform
 - Ansible
 - CloudFormation
 - Code Modularity and Re-Use
 - States and Secret Management
 - Provisioning Infrastructure using IAC

What is the agenda?

- Why do we need yet another framework (IAC)?
- What is Terraform and its benefits
- Terraform Basics
 - Installation
 - Commands, Workflows, Resource creation, File structure
- Some (not all) Best Practices
- Additional Resources for Learning and Practice

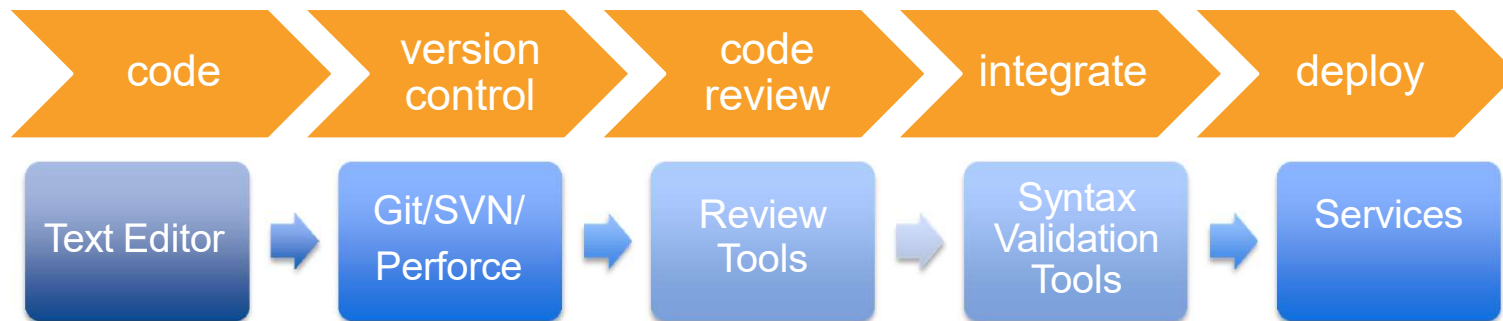
IAC – Why today?

- Moving to cloud based infrastructure opens doors to building and managing infrastructure in completely new ways:
- Infrastructure can be provisioned in seconds
- Scale can be achieved without complicated capacity planning
- Being API driven means we can interact with our infrastructure via languages typically used in our applications

Infrastructure as Code

Infrastructure as Code is a practice by where traditional infrastructure management techniques are supplemented and often replaced by using code based tools and software development techniques.

Infrastructure as Code Workflow



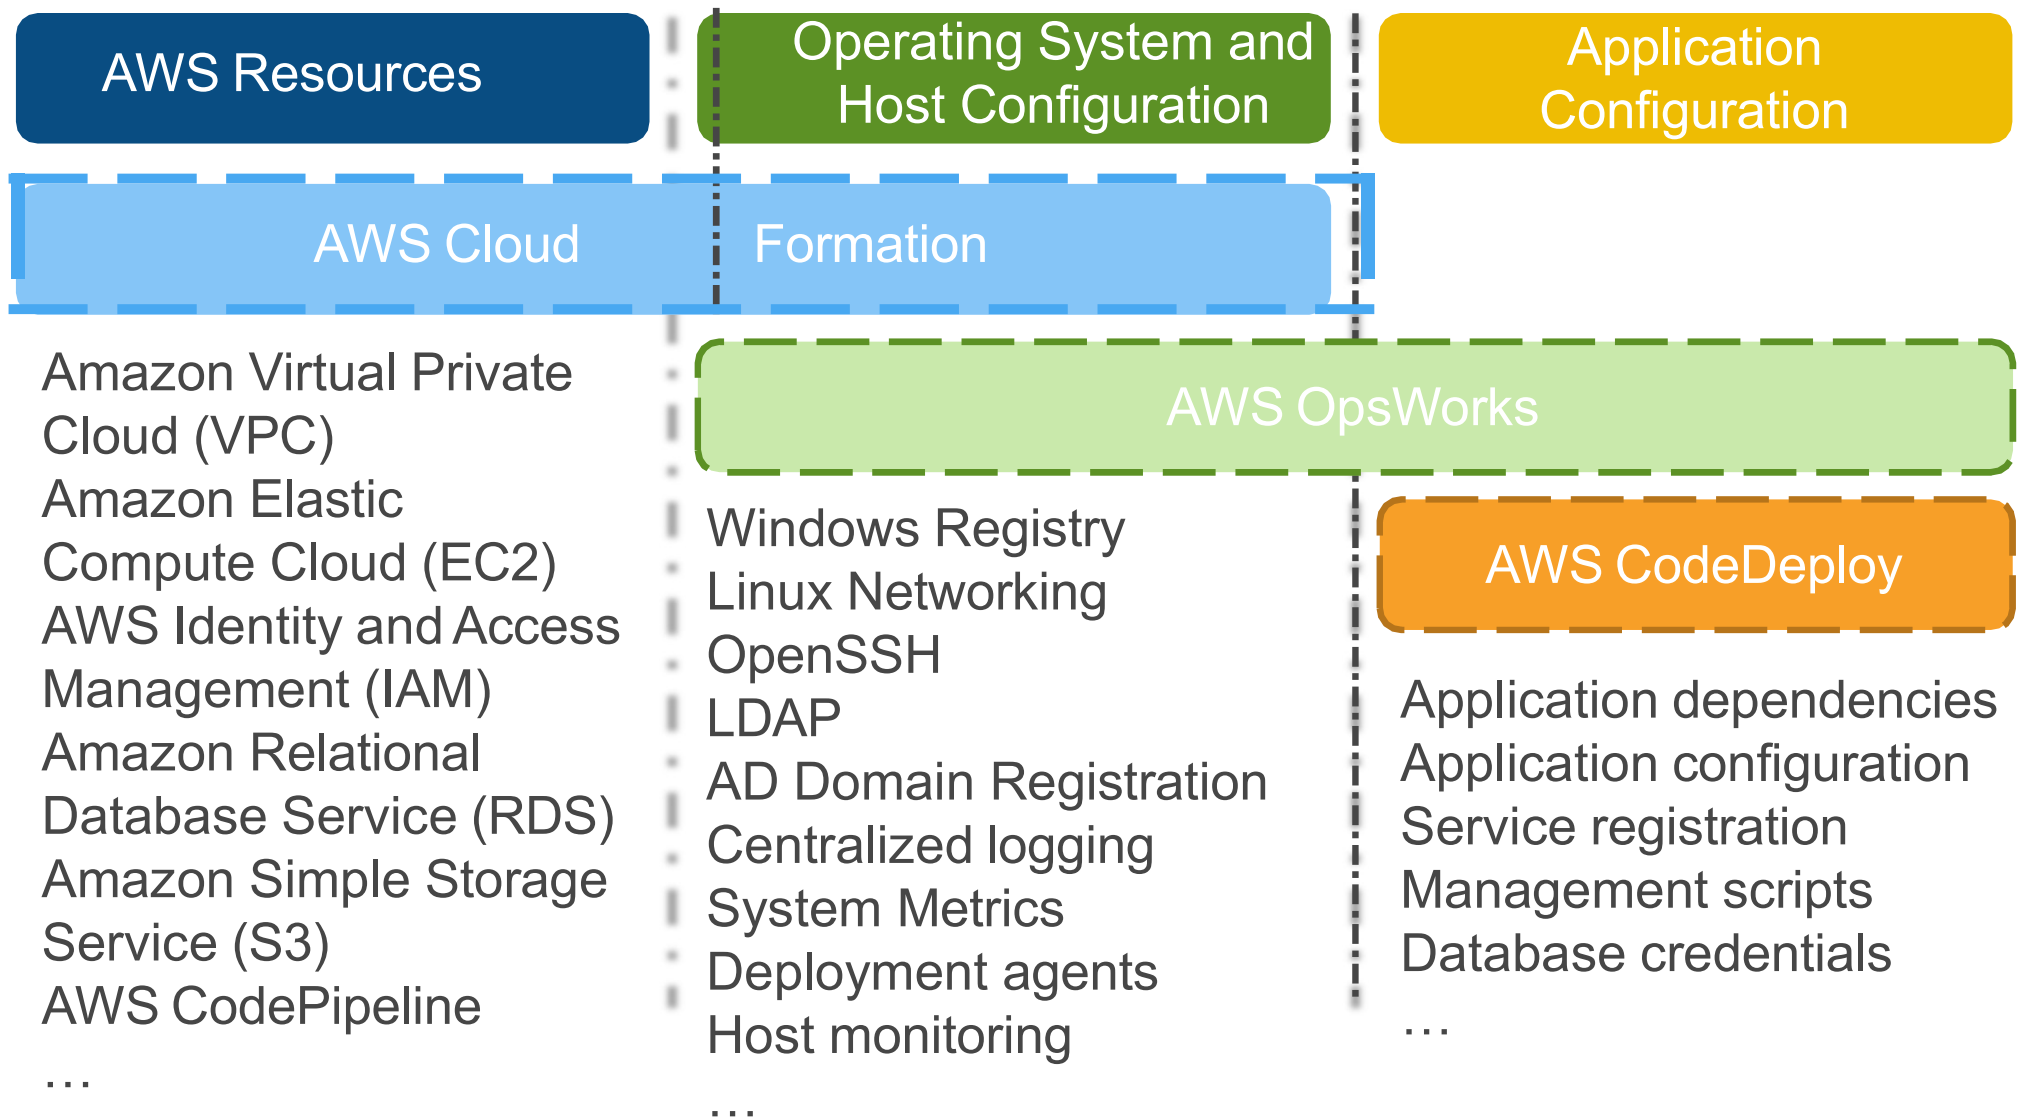
“It’s all software”

IT IS **ALL** CODE

Application Configuration

Operating System and Host Configuration

Resources



AWS CloudFormation

- AWS CloudFormation – more official quickstarts available
- Supports both JSON and YAML
- Need to use nested Templates to achieve re-usable (gone away since 2020 😊 in form of modules)
- So why not AWS CloudFormation directly?
 - Terraform is cloud Agnostic. The earlier slide could be done for Azure also
 - Dynamically Create Resources (for_each, count, dynamic blocks)
 - Built-in functions



AWS
CloudFormation

Create templates of your infrastructure

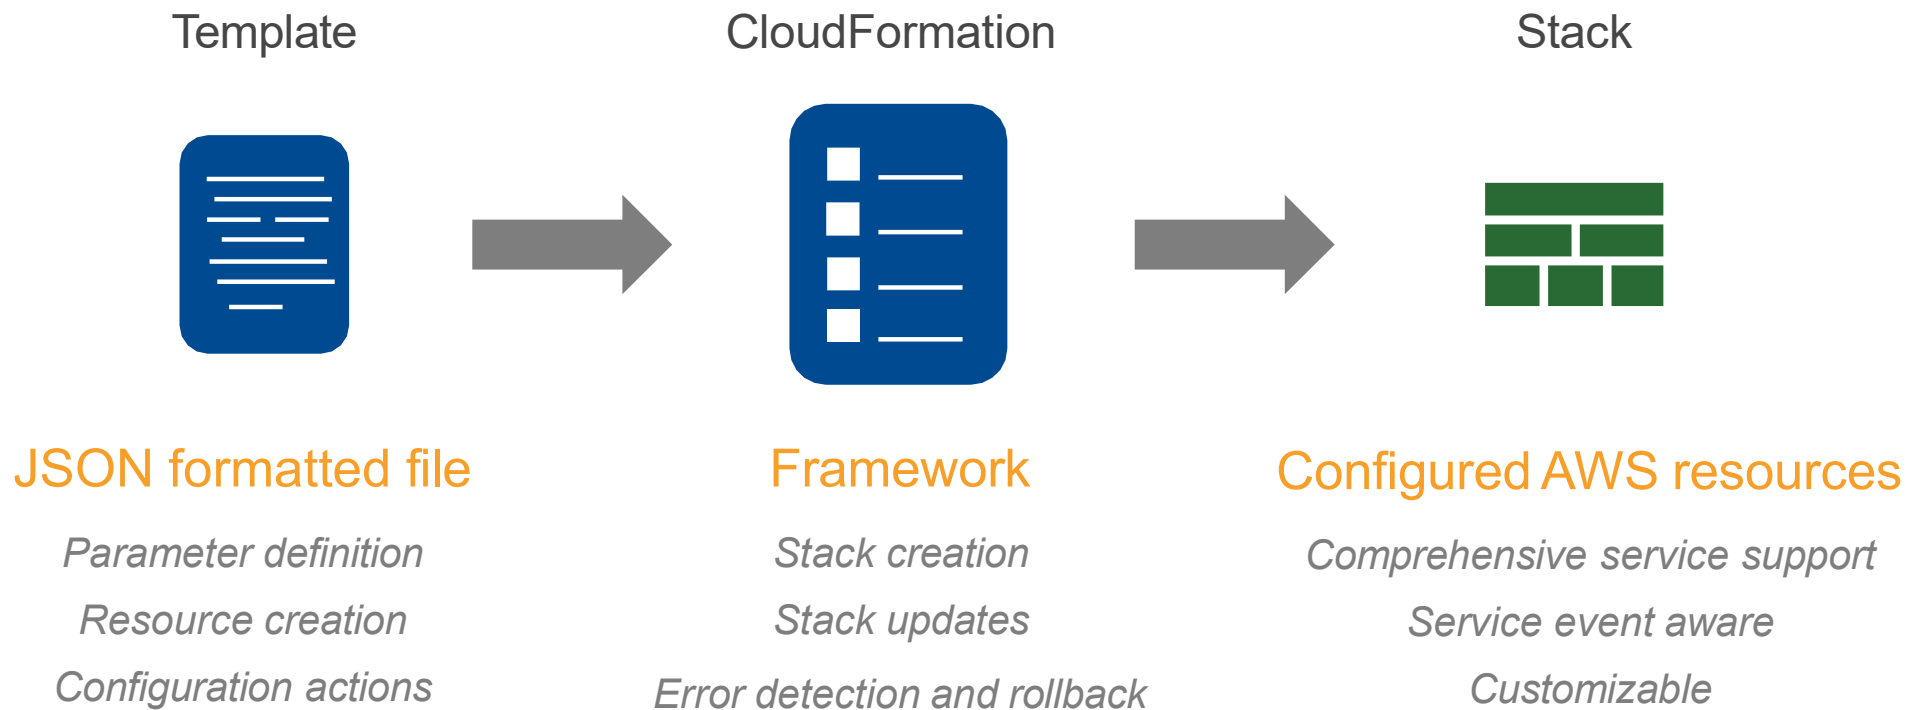
CloudFormation provisions AWS resources based on dependency needs

Version control/replicate/update templates like code

Integrates with development, CI/CD, management tools

Launched in 2010

CloudFormation – Components & Technology



Plain text

*Perfect for
version control*

JSON

Validatable

```
{
  "AWSTemplateFormatVersion" : "2010-09-09",

  "Description": "AWS CloudFormation Sample Template EC2InstanceSample: **WARNING** This template an Amazon EC2 instances. You will be billed for the AWS resources used if you create a stack from this template.",

  "Parameters" : {
    "KeyName" : {
      "Description": "Name of an existing EC2 KeyPair to enable SSH access to the instance",
      "Type" : "String"
    },
    "Environment": {
      "Type" : "String",
      "Default" : "Dev",
      "AllowedValues" : ["Mgmt", "Dev", "Staging", "Prod"],
      "Description" : "Environment that the instances will run in."
    }
  },

  "Mappings" : {
    "RegionMap" : {
      "us-east-1" : { "AMI" : "ami-7f418316" },
      "us-west-2" : { "AMI" : "ami-16fd7026" }
    }
  },

  "Conditions" : {
    "EnableEBSOptimized" : { "Fn::Equals" : [{"Ref" : " Environment "}, "Prod"]},
  },

  "Resources" : {
    "Ec2Instance" : {
      "Type" : "AWS::EC2::Instance",
      "Properties" : {
        "KeyName" : { "Ref" : "KeyName" },
        "EbsOptimized " : { "Fn::If": [ " EnableEBSOptimized ", {"true"}, {"false"} ]},
        "ImageId" : { "Fn::FindInMap" : [ "RegionMap", { "Ref" : "AWS::Region"}, "AMI" ]},
        "UserData" : { "Fn::Base64" : "80" }
      }
    },
  },

  "Outputs" : {
    "InstanceId" : {
      "Description" : "InstanceId of the newly created EC2 instance",
      "Value" : { "Ref" : "Ec2Instance" }
    },
    "PublicDNS" : {
      "Description" : "Public DNSName of the newly created EC2 instance",
      "Value" : { "Fn::GetAtt" : [ "Ec2Instance", "PublicDnsName" ] }
    }
  }
}
```

```
{
  "AWSTemplateFormatVersion": "2010-09-09",

  "Description": "AWS CloudFormation Sample Template EC2InstanceSample: **WARNING** This template an Amazon EC2 instances. You will be billed for the AWS resources used if you create a stack from this template.",

  "Parameters": {
    "KeyName": {
      "Description": "Name of an existing EC2 KeyPair to enable SSH access to the instance",
      "Type": "String"
    },
    "Environment": {
      "Type": "String",
      "Default": "Dev",
      "AllowedValues": ["Mgmt", "Dev", "Staging", "Prod"],
      "Description": "Environment that the instances will run in."
    }
  },

  "Mappings": {
    "RegionMap": {
      "us-east-1": { "AMI": "ami-7f418316" },
      "us-west-2": { "AMI": "ami-16fd7026" }
    }
  },

  "Conditions": {
    "EnableEBSOptimized": { "Fn::Equals": [{"Ref": "Environment"}, "Prod"] }
  },

  "Resources": {
    "Ec2Instance": {
      "Type": "AWS::EC2::Instance",
      "Properties": {
        "KeyName": { "Ref": "KeyName" },
        "EbsOptimized": { "Fn::If": [ "EnableEBSOptimized", { "true": true }, { "false": false } ] },
        "ImageId": { "Fn::FindInMap": [ "RegionMap", { "Ref": "AWS::Region" }, "AMI" ] },
        "UserData": { "Fn::Base64": "80" }
      }
    }
  },

  "Outputs": {
    "InstanceID": {
      "Description": "InstanceID of the newly created EC2 instance",
      "Value": { "Ref": "Ec2Instance" }
    },
    "PublicDNS": {
      "Description": "Public DNSName of the newly created EC2 instance",
      "Value": { "Fn::GetAtt": [ "Ec2Instance", "PublicDnsName" ] }
    }
  }
}
```

HEADERS

PARAMETERS

MAPPINGS

CONDITIONALS

RESOURCES

OUTPUTS

Copyrightt @2025 Seshagiri Sriram

Description of what your stack does, contains, etc

HEADERS

Provision time values that add structured flexibility and customization

PARAMETERS

Pre-defined conditional case statements

MAPPINGS

Conditional values set via evaluations of passed references

CONDITIONALS

AWS resource definitions

RESOURCES

Resulting attributes of stack resource creation

OUTPUTS

AWS CloudFormation

Templates (in action):

“AWSRegionVirt2AMI” Map

AWS::Region Pseudo
Parameter

```
"ImageId" : { "Fn::FindInMap" : [ "AWSRegionVirt2AMI", { "Ref" : "AWS::Region" },  
{"Fn::FindInMap": ["AWSInstanceType2Virt", { "Ref" : "myInstanceType" }, "Virt"]} ]},
```

“AWSInstanceType2Virt” Map

“myInstanceType”
Parameter

Mappings

Parameters:

```
"myInstanceType" : {  
  "Type" : "String",  
  "Default" : "t2.large",  
  "AllowedValues" : ["t2.micro", "t2.small",  
    "t2.medium", "t2.large"],  
  "Description" : "Instance type for instances created,  
    must be in the t2 family."  
}
```

Mappings:

```
"AWSInstanceType2Virt": {  
  "t2.micro": {"Virt": "HVM"},  
  "t2.small": {"Virt": "HVM"},  
  "t2.medium": {"Virt": "HVM"},  
  "t2.large": {"Virt": "HVM"},  
}
```

Mappings:

```
"AWSRegionVirt2AMI": {  
  "us-east-1": {  
    "PVM": "ami-50842d38",  
    "HVM": "ami-08842d60"  
  },  
  "us-west-2": {  
    "PVM": "ami-af86c69f",  
    "HVM": "ami-8786c6b7"  
  },  
  "us-west-1": {  
    "PVM": "ami-c7a8a182",  
    "HVM": "ami-cfa8a18a"  
  },  
}
```

Bootstrapping Applications


Option 1: Use EC2 UserData, which is available as a property of `AWS::EC2::Instance` resources

```
"Resources" : { "Ec2Instance"
  : {
    "Type" : "AWS::EC2::Instance",
    "Properties" : {
      "KeyName" : { "Ref" : "KeyName" },
      "SecurityGroups" : [ { "Ref" : "InstanceSecurityGroup" } ],
      "ImageId" : { "Fn::FindInMap" : [ "RegionMap", { "Ref" : "AWS::Region" }, "AMI" ] },
      "UserData" : { "Fn::Base64" : { "Fn::Join" : [ "", [
        "#!/bin/bash -ex","\n",
        "yum -y install gcc-c++ make","\n",
        "yum -y install mysql-devel sqlite-devel","\n",
        "yum -y install ruby-rdoc rubygems ruby-mysql ruby-devel","\n", "gem
        install --no-ri --no-rdoc rails","\n",
        "gem install --no-ri --no-rdoc mysql","\n", "gem
        install --no-ri --no-rdoc sqlite3","\n", "rails new
        myapp","\n",
        "cd myapp","\n",
        "rails server -d","\n" ] ] ] }
    }
  }
}
```

Copyrightt @2025 Seshagiri Sriram

Bootstrapping Applications & Handling Updates

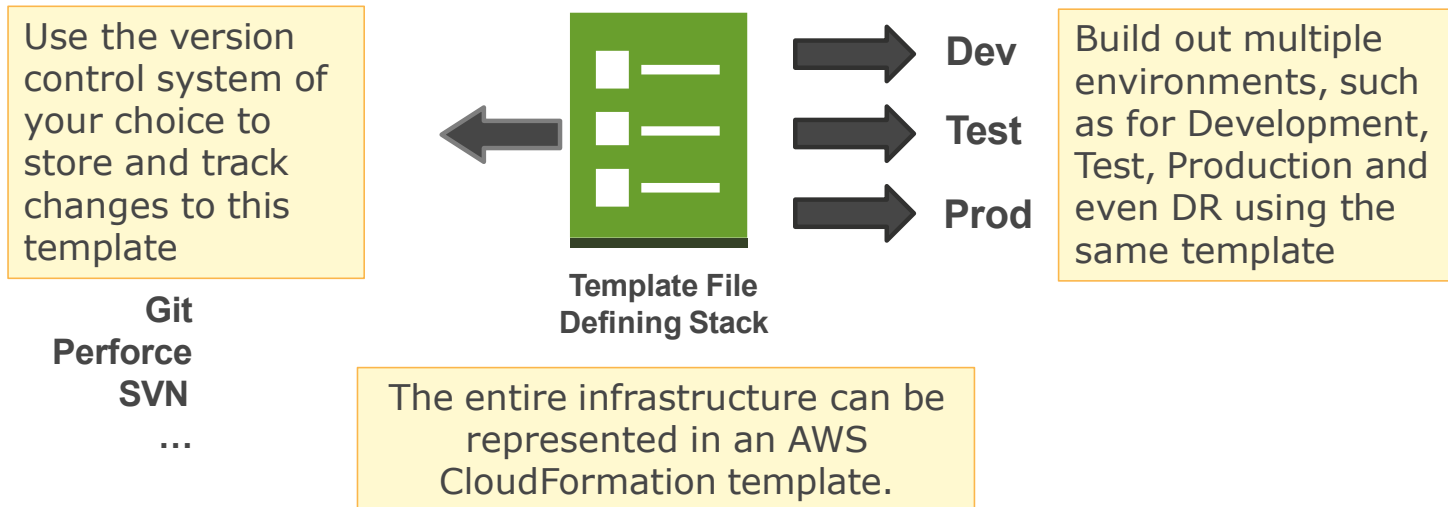
Option 2: AWS CloudFormation provides helper scripts for deployment within your EC2 instances



Metadata Key — `AWS::CloudFormation::Init`

Cfn-init reads this metadata key and installs the packages listed in this key (e.g., httpd, mysql, and php). Cfn-init also retrieves and expands files listed as sources.

Multiple Stacks/Environment



AWS CloudFormation

- **Versioning!**

- You track changes within your code Do it within your infrastructure!
 - What is changing?
 - Who made that change?
 - When was it made?
 - Why was it made?(tied to ticket/bug/project systems?)

- **Self imposed, but you need to be doing this!**

AWS CloudFormation

- **Testing:**

- Validate via API/CLI
 - `$ aws cloudformation validate-template` – confirm CF syntax
 - Use something like Jsonlint (<http://jsonlint.com/>) to find JSON issues like missing commas, brackets!
- Throw this into your testing/continuous integration pipelines!

AWS CloudFormation

- **Deploy & update via console or API/command line:**
 - Just a couple of clicks OR
 - `aws cloudformation create-stack --stack-name myteststack --template-body file:///home//local//test//sampletemplate.json --parameters ParameterKey=string,ParameterValue=string`

What happens once deployed?

Ongoing Management

- Updates/patches?
- New software?
- New configurations?
- New code deploys?
- Pool specific changes?
- Environment specific changes?
- Run commands across all hosts?
- Be on top of all running resources?


Select Right Tool

- Can be done with AWS CloudFormation – maintenance will be harder.
- Consider AWS OpsWork and Chef

Infrastructure-as-Code (IAC) Benefits

- ✓ Reproducible Environments
- ✓ Automation – CI/CD
- ✓ Trackable – GitHub
- ✓ Language – HCL (Hashicorp Configuration Language)
- ✓ Workflow
- ✓ Multiple Providers
- ✓ Can provide similar configuration across multiple cloud vendors
- ✓ Similar infra across environments

What is Terraform?



- A templating language HashiCorp Configuration Language (**HCL**)
- A tool that can be used to orchestrate the provisioning of:
 - Public clouds (Azure, AWS, GCP, Oracle, Alibaba)
 - On-premises (VMware)
 - Other (Cisco, GitHub, GitLab, New Relic, Okta, Rabbit MQ)
- Uses State files (more on this later)
- Is **NOT** used for configuration
 - PowerShell Desired State Configuration (DSC), Chef, Puppet, Ansible

TERRAFORM <> ANSIBLE

- TERRAFORM – IAC
- ANSIBLE – CONFIGURATION MANAGEMENT
- ANSIBLE
 - Makes configuration management and application deployment easier
- TERRAFORM
 - building, updating, and versioning infrastructure safely and effectively

Terraform

- Clean and easy code to write and maintain
- Fully declarative configuration
- Version control on infrastructure
- Implicit dependencies management – explicit can be forced
- Ecosystem of providers and skilled personnel

Terraform

- Multiple Cloud Providers
- Declarative Configuration
- Execution Plans
- Resource Graphs
- State Management
- Loops, Conditionals, Dynamic Blocks
-

Terraform




Copyright @ 2020 Seshagiri Sriram

Getting Started

- Installation on Windows is pretty straightforward
- Download AMD64 version for Windows from [here](#).
- Unzip the file and copy terraform to a folder that exists in PATH

How does it work?

- Terraform **init**
 - Initializes the current working directory
- Terraform **plan**
 - Execution plan to validate against existing environment
- Terraform **apply**
 - Deploys and updates resources
- Terraform **destroy**
 - Removes all resources defined in a configuration



Resources

- Resources (e.g. an NGINX server) have
 - Type
 - A **required Provider**
 - Name (internal to Terraform)
 - Configuration with Deployment Details

```
resource "azurerm_resource_group" "SharedServicesRG" {  
  name      = "SharedServicesRG"  
  location = "Canada Central"  
}
```

Resource Type Name

Resource Configuration

Backends

- Determines how state is loaded and how an operation such as **apply** is executed
- By default, Terraform uses the "local" backend
- Benefits of backends:
 - Working in a team
 - Store state remotely and protect state with locks to prevent corruption
 - Keeping sensitive information off disk
 - Retrieved on demand and only stored in memory
 - Remote operations
 - Remote execution

State

- State keeps track of the all managed resources and their associated properties with current values.
- **Essential for managing changes to infrastructure over time**
- Preserve the state file for the entire life cycle of the resources
 - You can create a separate state file per resource group, application, shared service (ie. core networking)
 - Terraform Workspaces should also be used to separate application and environment boundaries
- Recommended to use a **remote backend** to save state in centralized, secure storage
 - Example: Storage account, Terraform Cloud, Terraform Enterprise, Artifactory, Consul
- You must initialize the Terraform State
 - This is what **terraform init** does

State

**Secrets (like
usernames/passwords,
access keys/tokens, etc.)
can be written to your state
file!**

Providers

- The provider block is used to configure the named provider
- Is responsible for creating and managing resources, and for all other interactions including authentication
- The version argument is optional, but recommended

```
provider "azurerm" {  
  version = ">=2.0.0"  
  subscription_id = "<<REMOVED>>"  
  client_id       = "<<REMOVED>>"  
  client_secret   = "<<REMOVED>>"  
  tenant_id      = "<<REMOVED>>"  
}
```

Variables

- Parameterize the configurations
- If no value is assigned to a variable and the variable has a **default** key in its declaration, that value will be used for the variable
- Can be provided...
 - Within the Terraform template
 - Within its own Terraform template file
 - Within a .TFVARS files
 - Through command-line
 - Through Environment variables

Variables

```
variable "location" {  
  type      = string  
  default   = "westus"  
  description = "Specify a location see: az account list-locations -o table"  
}  
  
resource "azurerm_resource_group" "example" {  
  name = var.resource_group_name  
  location = var.location  
}
```

Variable Precedence

- a. Command Line
- b. From a File
- c. From Environment Variables
- d. UI Input

Dependencies

- *Implicit* dependencies, which Terraform and the provider determine automatically based on the configuration
- *Explicit* dependencies, which you define using the `depends_on` meta-argument

Dependencies

```
resource "azurerm_resource_group" "example" {  
  name = var.resource_group_name  
  location = var.location  
}  
  
resource "azurerm_storage_account" "example" {  
  name = "storageaccountname"  
  resource_group_name = azurerm_resource_group.example.name  
  location = azurerm_resource_group.example.location  
  account_tier = "Standard"  
  account_replication_type = "GRS"  
  tags = {  
    environment = "staging"  
  }  
}
```

Outputs

- Used to organize data to be easily queried and shown back to the Terraform user
- Data is outputted when **apply** is called
- Outputs can be queried after a run using the **terraform output <<output name>>** command

```
output "SharedServices-RGName" {  
  value = azurerm_virtual_network.SharedServicesVNET.*.resource_group_name  
}  
  
output "SharedServices-VNet-Name" {  
  value = azurerm_virtual_network.SharedServicesVNET.*.name  
}  
  
output "SharedServices-VNet-ID" {  
  value = azurerm_virtual_network.SharedServicesVNET.*.id  
}
```

Best Practices

- Use remote backends
- Manage Terraform, providers and modules versions
- Use implicit dependencies
- Use modules (custom or from the HashiCorp public registry <https://registry.terraform.io>)
- Use ARM templates only if you don't have another choice

Terraform – Adv. State and Secrets

1.Pre-requisite #1: Don't Store Secrets in Plain Text

2.Pre-requisite #2: Keep Your Terraform State Secure

1.Technique #1: Environment Variables

2.Technique #2: Encrypted Files (e.g., KMS, PGP, SOPS)

3.Technique #3: Secret Stores (e.g., Vault, AWS Secrets manager)

Terraform – Adv. State and Secrets

- Method #1: Terraform has native support for Environment Variables

```
variable "username" {  
  description = "The username for the DB master user"  
  type        = string  
  sensitive   = true  
}  
variable "password" {  
  description = "The password for the DB master user"  
  type        = string  
  sensitive   = true  
}
```

```
resource "aws_db_instance" "example" {  
  engine           = "mysql"  
  engine_version   = "5.7"  
  instance_class   = "db.t2.micro"  
  name             = "example" # Set the secrets from  
  variables  
  username         = var.username  
  password         = var.password  
}
```

```
export TF_VAR_username='someusername'  
export TF_VAR_password='Somepassword'  
# neither user name nor password will be found in Version Control
```

Terraform – Adv. State and Secrets

- Method #2: Use tools for storing sensitive secrets e.g pass

```
$ pass insert db_username  
Enter password for db_user:  
$ pass insert db_password  
Enter password for db_password
```

- You still have to remember the password 😊.
- Secrets can be read as:

```
$ pass db_username
```

```
export TF_VAR_username=$(pass db_username)  
export TF_VAR_password=$(pass db_password)  
# neither user name nor password will be found in Version Control
```


Terraform – Adv. State and Secrets

- Use either
 - AWS KMS
 - Azure Key Vault
 - GCP KMS
- Basic process consists of creating a YML file like below (This file **SHOULD NOT BE** version controlled)

```
username: someuser  
password: somepassword
```

Terraform – Adv. State and Secrets

- Encrypt the file using the corresponding AWS/GCP/Azure KMS Key
- Check in the encrypted file into version control.
- Use the encrypted file like so:

```
data "aws_kms_secrets" "creds" {  
  secret {  
    name = "db"  
    payload = file("${path.module}/db-creds.yml.encrypted")  
  }  
}  
  
locals {  
  db_creds = yamldecode(data.aws_kms_secrets.creds.plaintext["db"])  
}  
...  
5 Username = locals.db_creds.username
```

Terraform – Adv. State and Secrets

- Use SOPS and Terragrunt (out of scope) to simplify the workflow of editing and working with secret files.

Modular Code

- A *module* is a container for multiple resources that are used together.
- You can use modules to create lightweight abstractions, so that you can describe your infrastructure in terms of its architecture, rather than directly in terms of physical objects.
- The .tf files in your working directory when you run terraform plan or terraform apply together form the root module.
- That module may call other modules and connect them together by passing output values from one to input values of another.

Modular Code – Structure

- Modules use
 - Input Variables
 - Accept Values from calling module
 - Output Values
 - Send Values to calling module
 - Resources
 - Define one or more infra objects
- Each Module is in a separate directory
- Modules can be loaded locally or from a remote repository

Modular Code – Structure

- Modules can call other modules using module block
- Examples of a module

```
.  
├── main.tf  
└── terraform-azurevm-vnet  
    ├── main.tf  
    ├── output.tf  
    └── variables.tf
```

Modular Code – Structure

```
module "vnet" {  
  source = "../terraform-azurerm-vnet"  
  resource_group_name = "test-rg"  
  location            = "westus"  
  vnet_name           = "new-vnet"  
  vnet_address_space  = "10.10.0.0/16"  
  subnet_name         = "subnet01"  
  subnet_address_space = "10.10.10.0/24"  
}
```

```
D:\02.Work\Engagements\demo_code\terraform\module_demo>terraform init
```

```
Initializing the backend...
```

```
Initializing modules...
```

```
Initializing provider plugins...
```

```
- Finding hashicorp/azurerm versions matching "3.0.0"...
```

```
- Installing hashicorp/azurerm v3.0.0...
```

```
- Installed hashicorp/azurerm v3.0.0 (signed by HashiCorp)
```

```
Terraform has created a lock file .terraform.lock.hcl to record the provider  
selections it made above. Include this file in your version control repository  
so that Terraform can guarantee to make the same selections by default when  
you run "terraform init" in the future.
```

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see any changes that are required for your infrastructure. All Terraform commands should now work.

If you ever set or change modules or backend configuration for Terraform, rerun this command to reinitialize your working directory. If you forget, other commands will detect it and remind you to do so if necessary.

Modular Code

- Now you can plan and apply.
- Note for this to work, you will need to install the AZ CLI and have a valid Azure Subscription.

Demo Time

Simple Terraform Deployments



We're done!
Thank you for your time and
participation.