

Kubernetes

Deploying Apps with Kubernetes

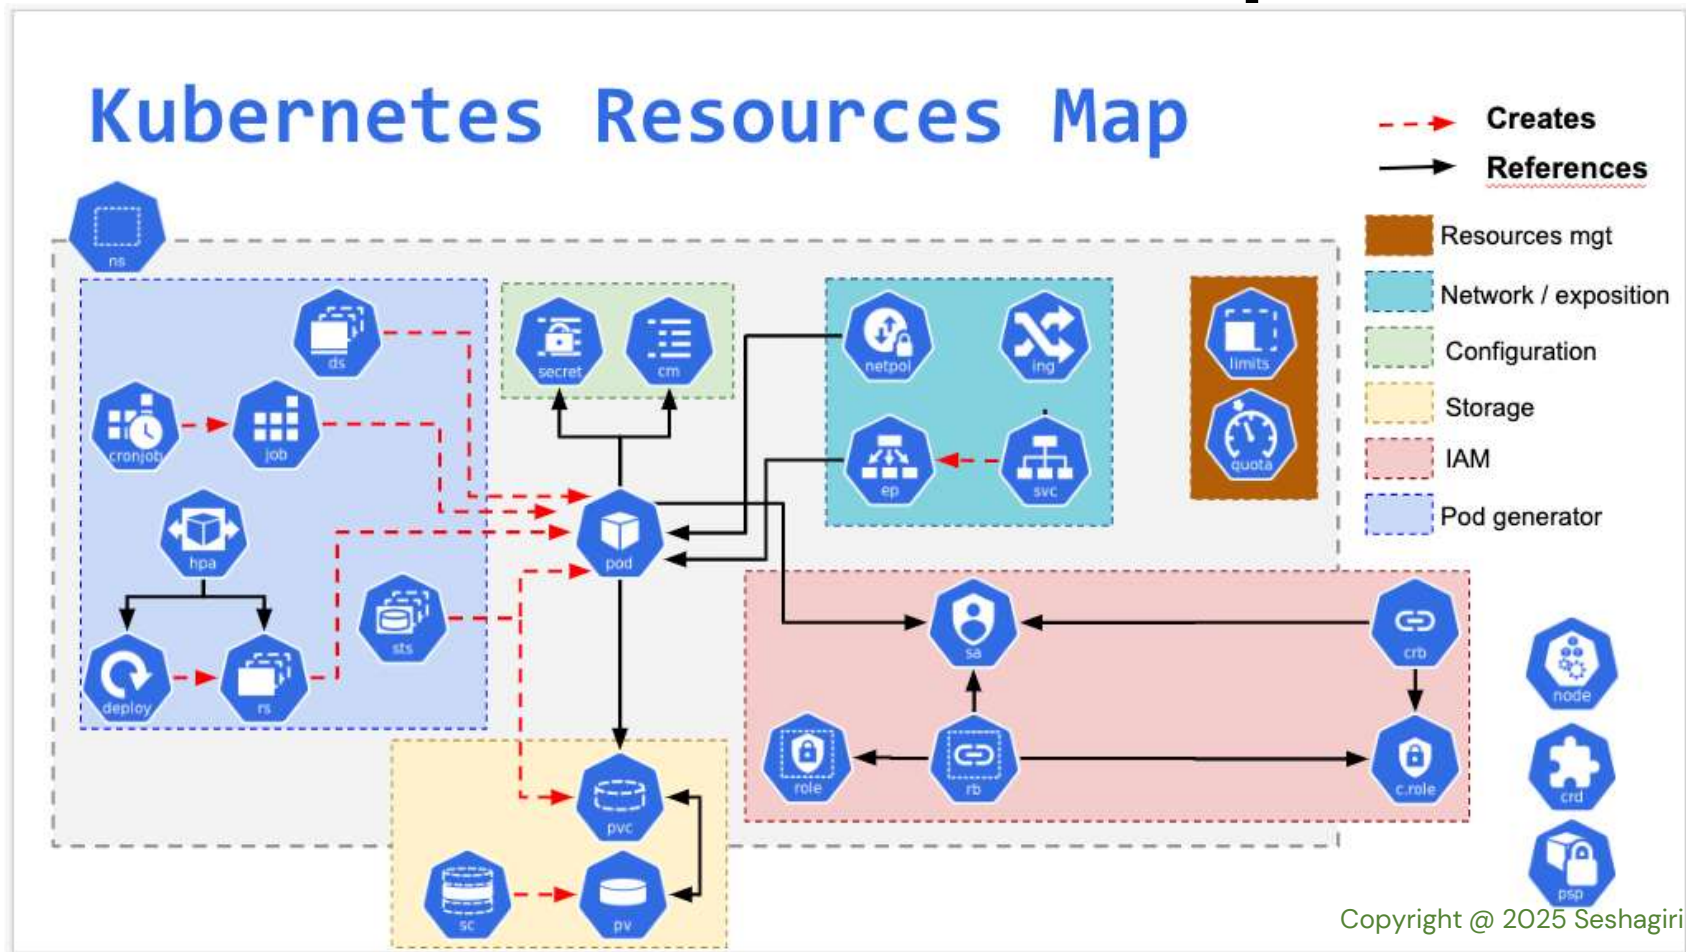
Kubernetes Introduction

- **What is Kubernetes (“k8s” or “kube”)**
- **Kubernetes Architecture**
- **Core Concepts of Kubernetes**
- Kubernetes resources explained
 - Networking
 - Storage and Persistence
 - Logging
- Application Optimization on Kubernetes
- Kubernetes effect on the software development life cycle
- Local and Distributed Abstractions and Primitives
- Container Design Patterns and best practices
- Deployment and release strategy with Kubernetes

Kubernetes Introduction



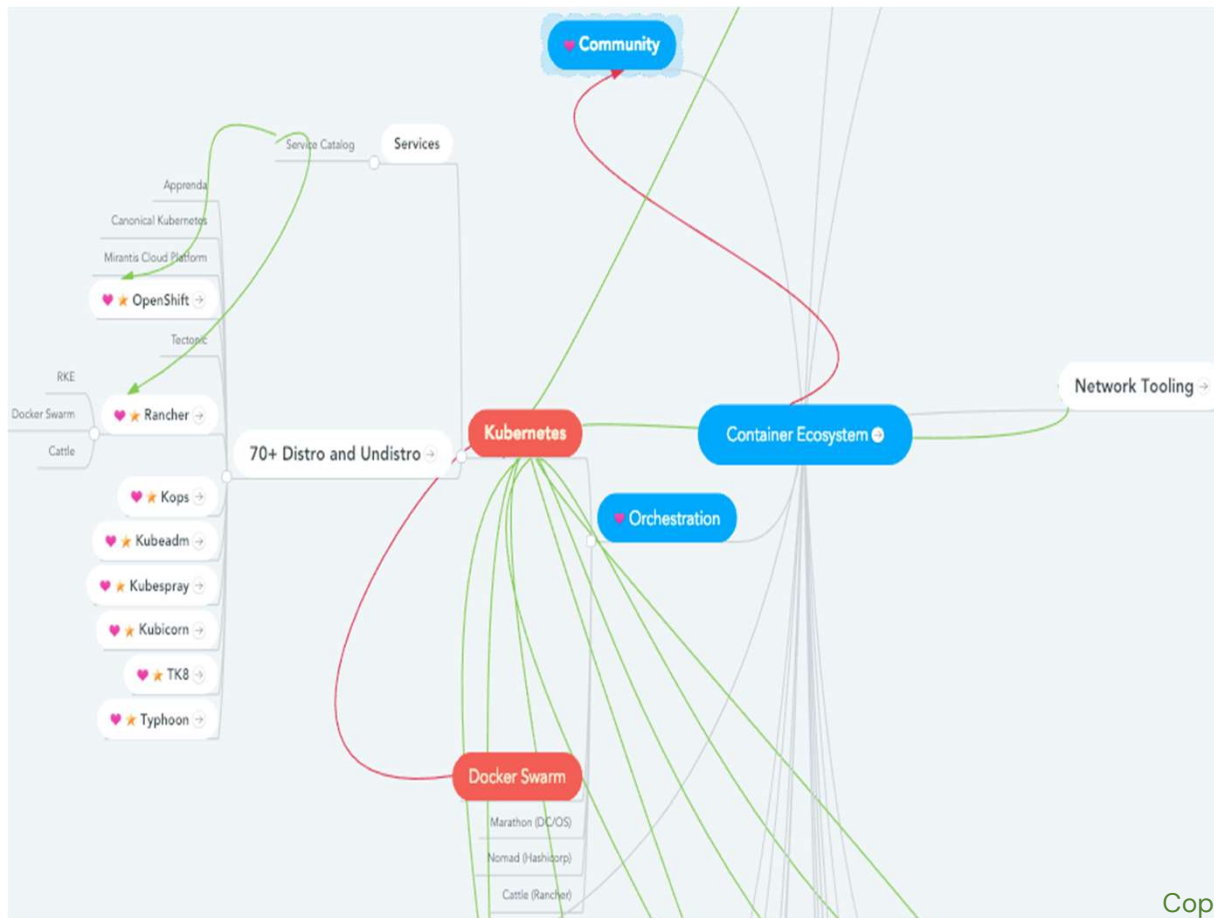
Kubernetes – Resource Map



Kubernetes Introduction

- What is Kubernetes
- Overview of Kubernetes Architecture
- Key Terms and Concepts
- Installation and Setup
- Exercises

Kubernetes Introduction



Kubernetes – What is it

- Kubernetes is Greek for "helmsman", your guide through unknown waters)
- Kubernetes is the linux kernel of distributed systems
- Kubernetes is the linux of the cloud!
- Kubernetes is a platform and container orchestration tool for automating deployment, scaling, and operations of application containers.
- Kubernetes supports, Containerd, CRI-O, Kata containers (formerly clear and hyper) and Virtlet

Kubernetes – Container Runtimes

- Kubernetes unlike docker does support multiple Container Runtimes
 - Docker
 - CRI-O
 - ContainerD
- The choice of the container runtime depends on a number of factors
- Additional Resources
 - Linux Container Internals by Scott McCarty
 - Container Runtimes and Kubernetes by Fahed Dorgaa
 - Kubernetes Runtime Comparison

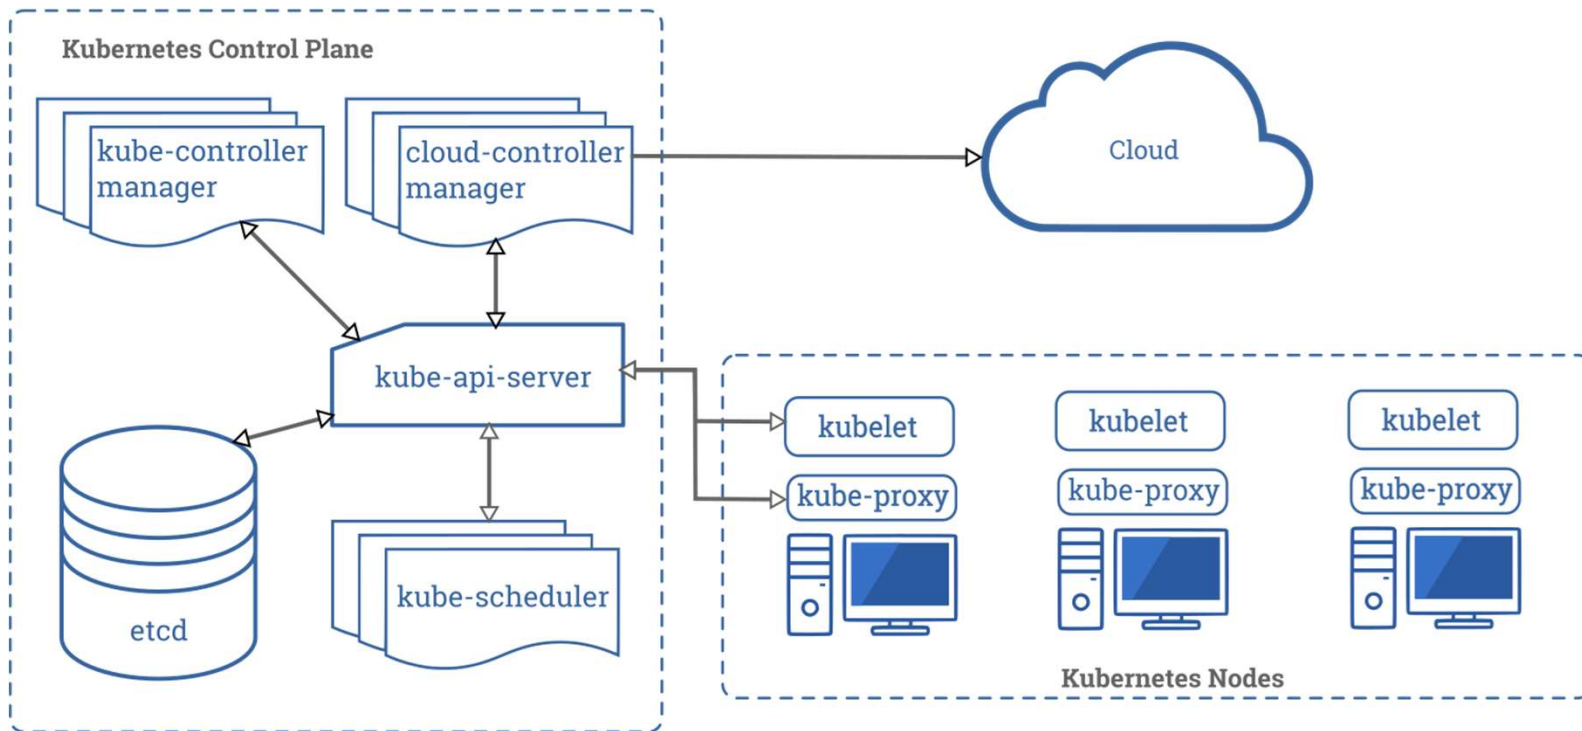
Kubernetes – Container Runtimes

- Kubernetes unlike docker does support multiple Container Runtimes
 - Docker
 - CRI-O
 - ContainerD
- The choice of the container runtime depends on a number of factors
- Additional Resources
 - Linux Container Internals by Scott McCarty
 - Container Runtimes and Kubernetes by Fahed Dorgaa
 - Kubernetes Runtime Comparison

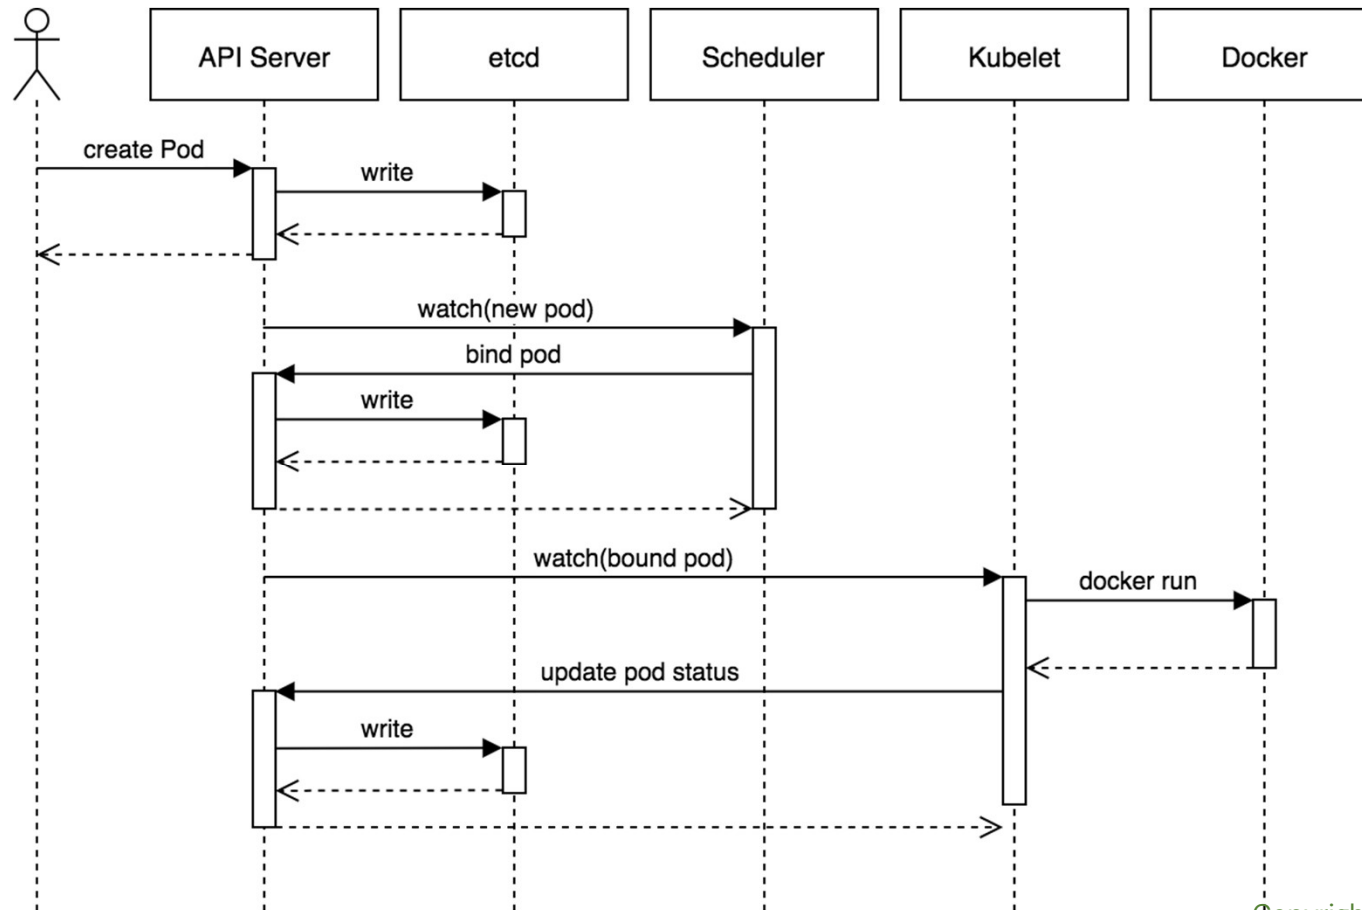
Kubernetes – How does it work

- A main Master Node
- Multiple Worker Nodes
 - NB: In MiniKube, only a single node
- Each Worker node can handle multiple pods
- Pods are just containers clustered as a single working unit
- We build applications and package them as pods using pod definitions and deployments to the main master node
- The master node takes these definitions and deploys to the worker nodes.

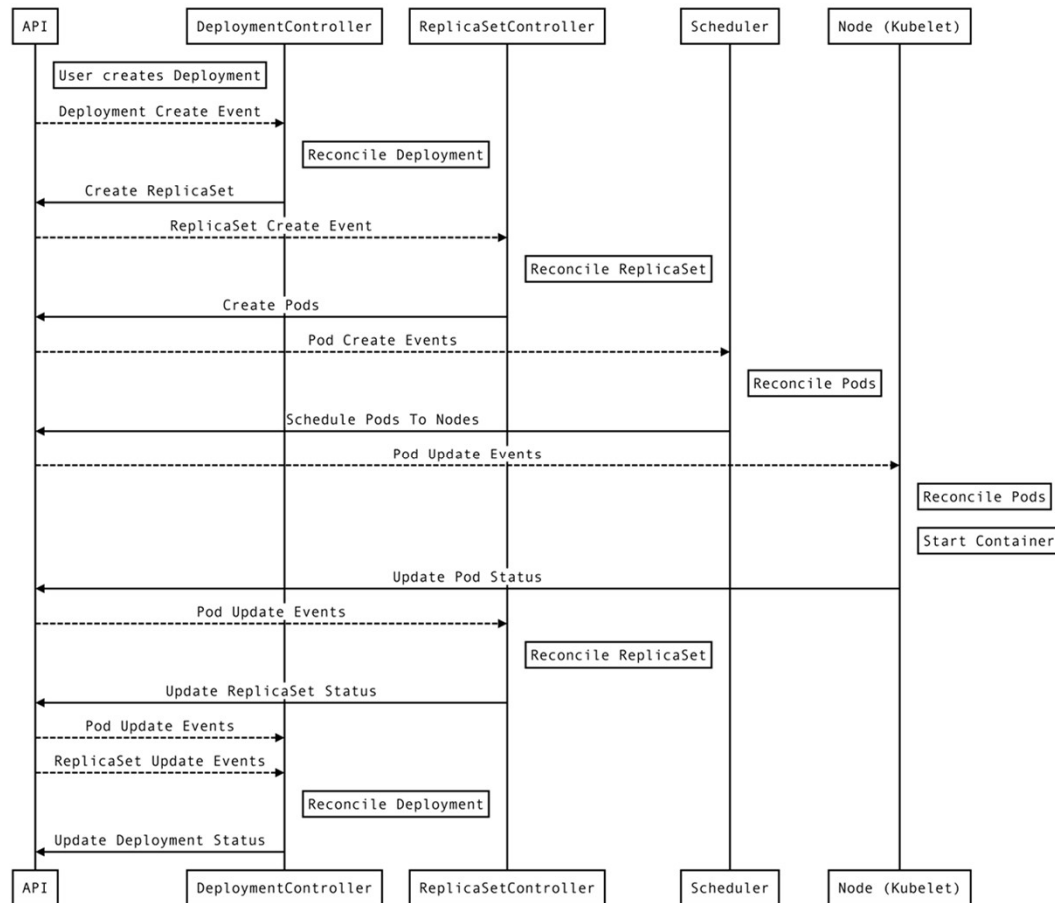
Kubernetes – Components



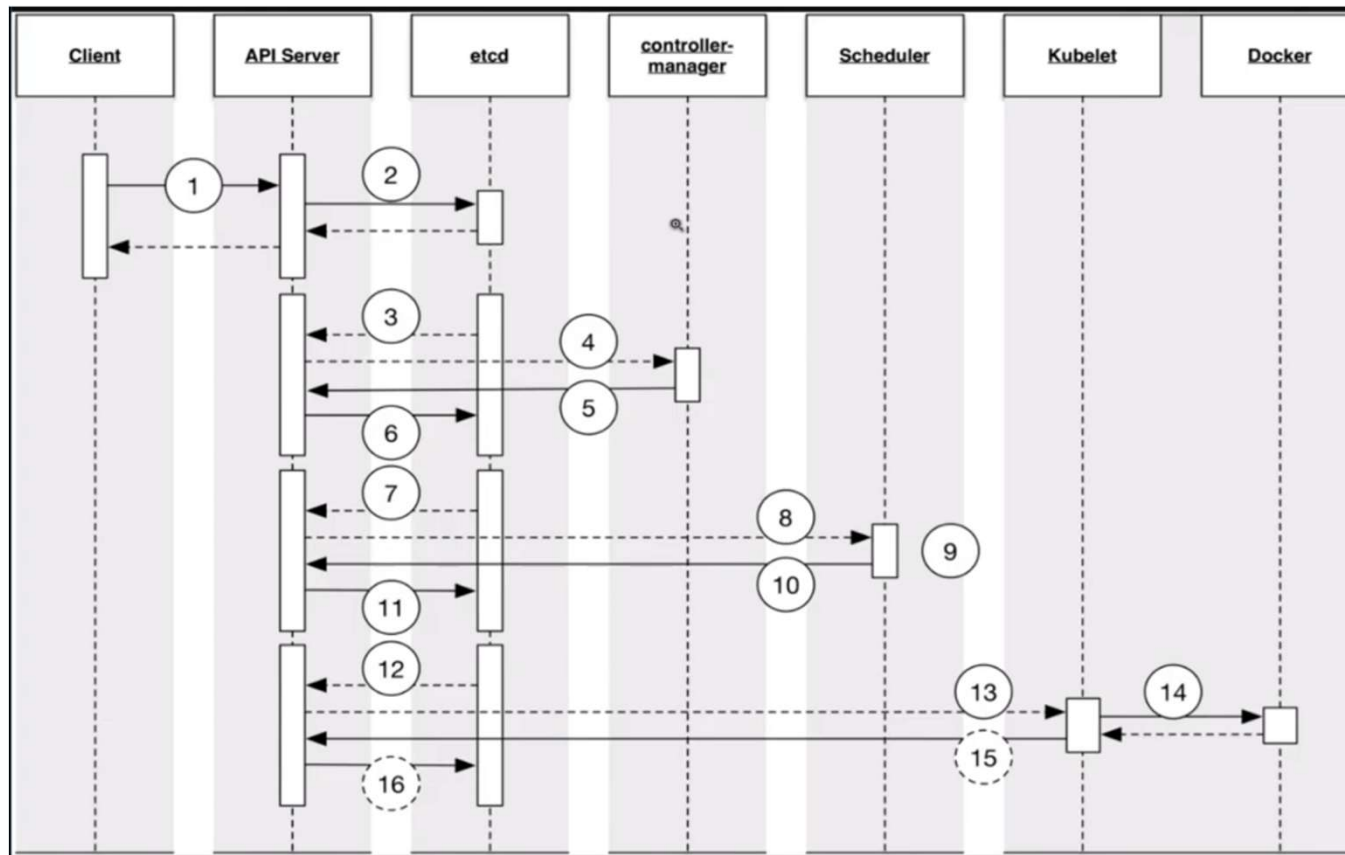
Kubernetes – The flow of work



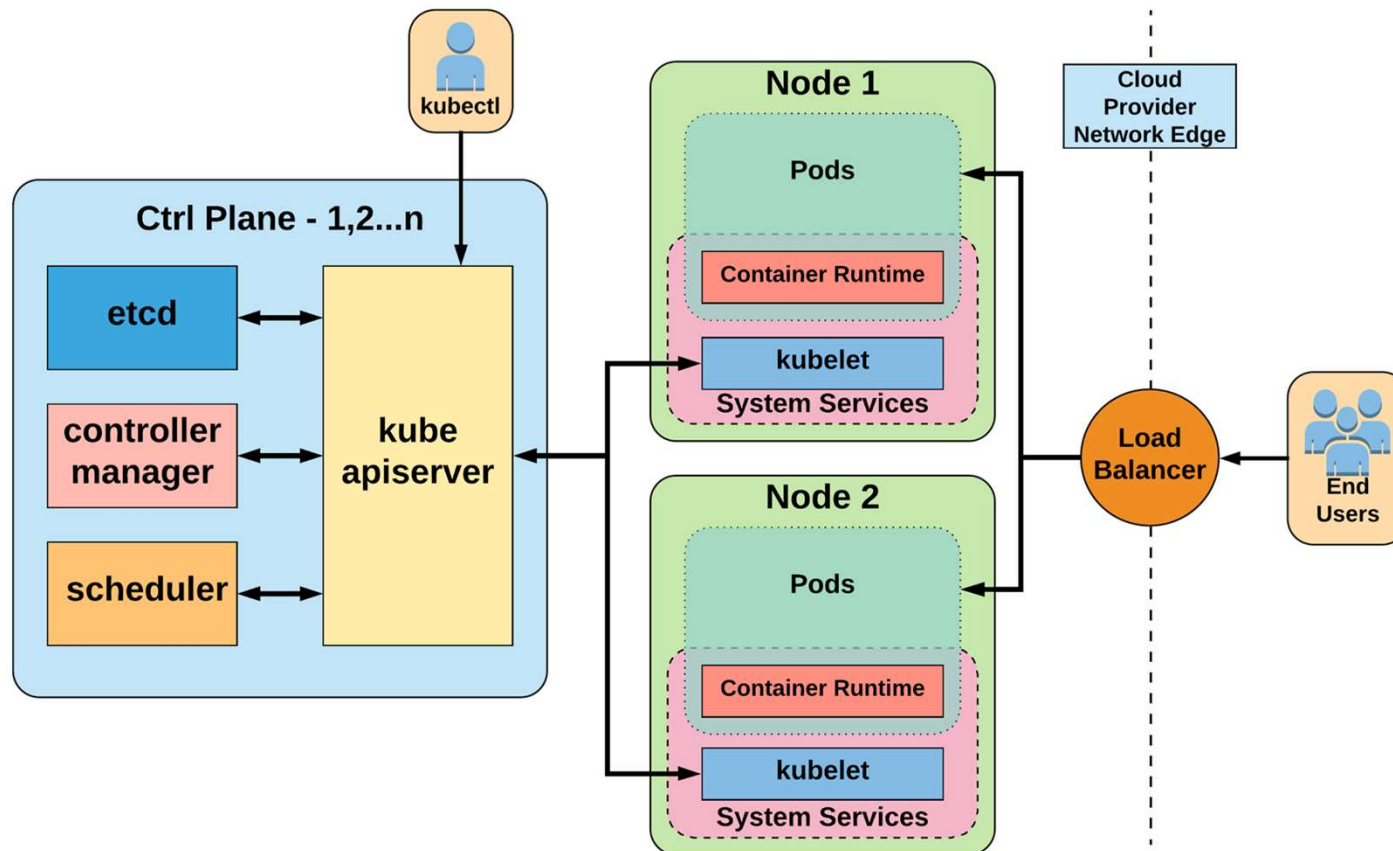
Kubernetes – The flow of work



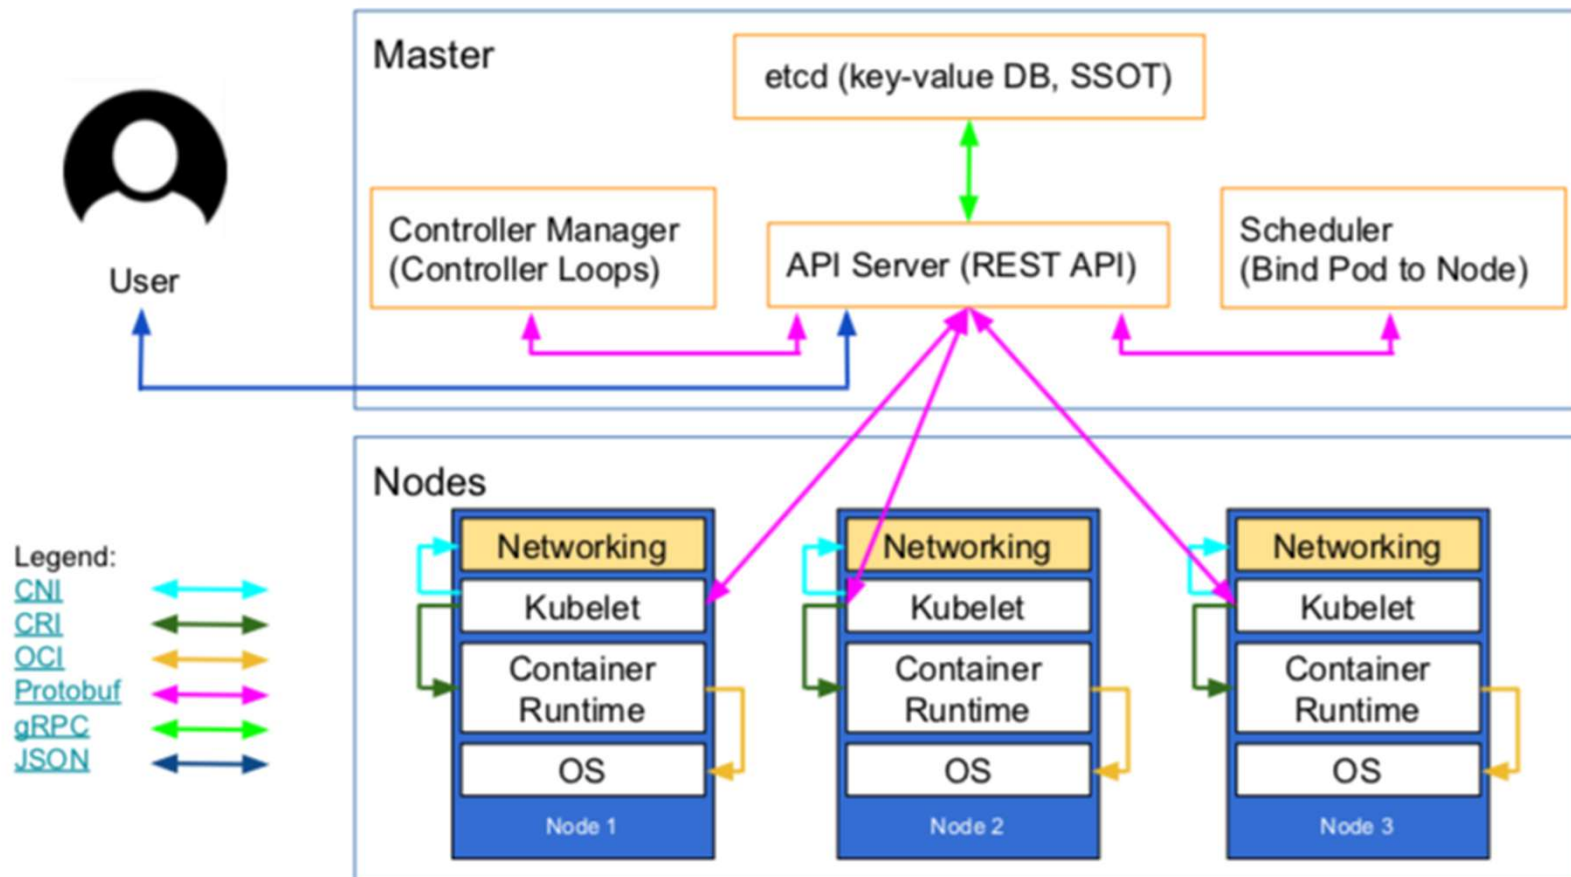
Kubernetes – The flow of work



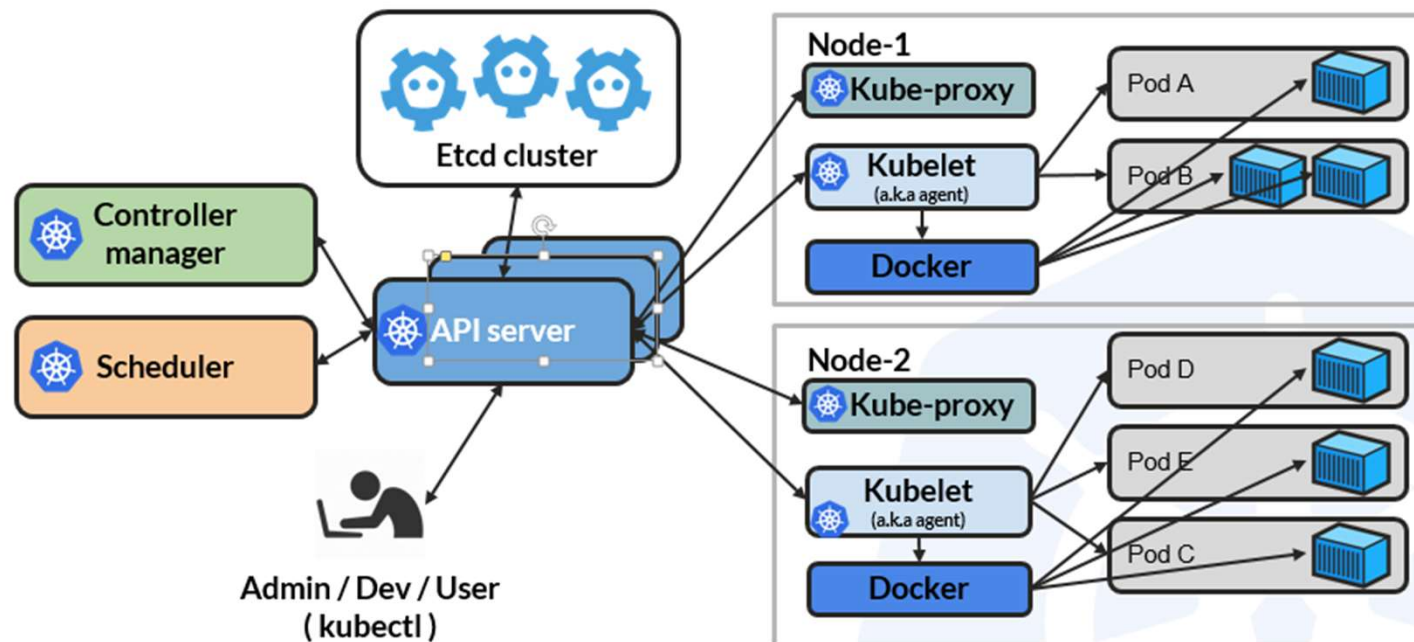
Kubernetes Architecture Overview



Kubernetes Architecture Overview



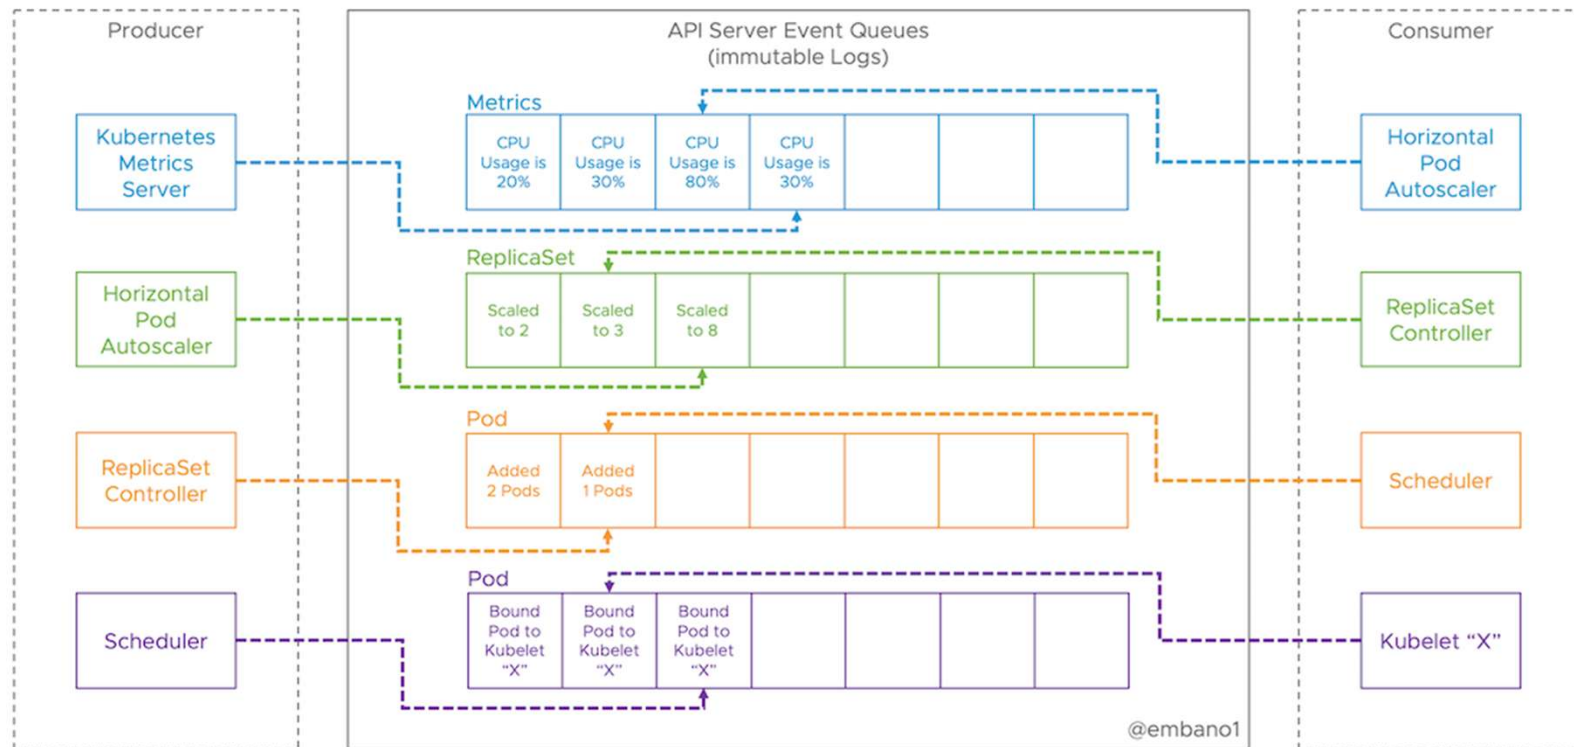
Kubernetes Architecture Overview



Kubernetes Architecture Overview

- Kubernetes is like Kafka
- Event Driven Architecture
- Producers write to a queue – this is considering as generating an event
- Consumers subscribe to these events and/or notified and respond accordingly

Kubernetes Architecture Overview



Kubernetes Core Concepts

- Pod
- Label and selectors
- Controllers
 - Deployments
 - ReplicaSet
 - ReplicationController
 - DaemonSet
- Service
- StatefulSets
- ConfigMaps
- Secrets
- Persistent Volumes (attaching storage to containers)
- Life Cycle of Applications in Kubernetes
 - Updating Pods
 - Rolling updates
 - Rollback

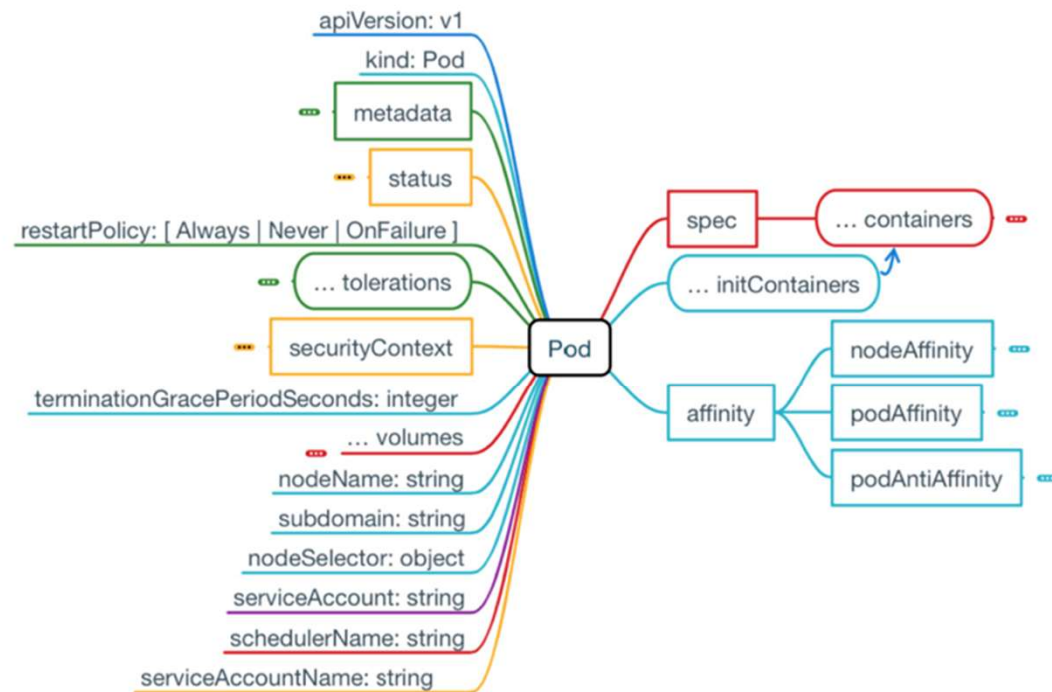
Kubernetes Core Concepts

	Resource (abbr.) [API version]	Description
	Namespace* (ns) [v1]	Enables organizing resources into non-overlapping groups (for example, per tenant)
Deploying Workloads	Pod (po) [v1]	The basic deployable unit containing one or more processes in co-located containers
	ReplicaSet	Keeps one or more pod replicas running
	ReplicationController	The older, less-powerful equivalent of a ReplicaSet

Kubernetes Core Concepts

	Resource (abbr.) [API version]	Description
Deploying Workloads	Job	Runs pods that perform a completable task
	CronJob	Runs a scheduled job once or periodically
	DaemonSet	Runs one pod replica per node (on all nodes or only on those matching a node selector)
	StatefulSet	Runs stateful pods with a stable identity
	Deployment	Declarative deployment and updates of pods

Kubernetes Core Concepts



Kubernetes Core Concepts

	Resource (abbr.) [API version]	Description
Services	Service (svc) [v1] Endpoints (ep) [v1] Ingress (ing) [extensions/v1beta1]	Exposes one or more pods at a single and stable IP address and port pair Defines which pods (or other servers) are exposed through a service Exposes one or more services to external clients through a single externally reachable IP address

Kubernetes Core Concepts

	Resource (abbr.) [API version]	Description
Config	ConfigMap (cm) [v1]	A key-value map for storing non-sensitive config options for apps and exposing it to them
	Secret [v1]	Like a ConfigMap, but for sensitive data
Storage	PersistentVolume* (pv) [v1]	Points to persistent storage that can be mounted into a pod through a PersistentVolumeClaim
	PersistentVolumeClaim (pvc) [v1]	A request for and claim to a PersistentVolume
	StorageClass* (sc) [storage.k8s.io/v1]	Defines the type of storage in a PersistentVolumeClaim

Kubernetes Core Concepts

	Resource (abbr.) [API version]	Description
Scaling	HorizontalPodAutoscaler (hpa) [autoscaling/v2beta1**]	Automatically scales number of pod replicas based on CPU usage or another metric
	PodDisruptionBudget (pdb) [policy/v1beta1]	Defines the minimum number of pods that must remain running when evacuating nodes
Resources	LimitRange (limits) [v1]	Defines the min, max, default limits, and default requests for pods in a namespace
	ResourceQuota (quota) [v1]	Defines the amount of computational resources available to pods in the namespace

Kubernetes Core Concepts

	Resource (abbr.) [API version]	Description
Cluster state	Node* (no) [v1]	Represents a Kubernetes worker node
	Cluster* [federation/v1beta1]	A Kubernetes cluster (used in cluster federation)
	ComponentStatus* (cs) [v1]	Status of a Control Plane component
	Event (ev) [v1]	A report of something that occurred in the cluster

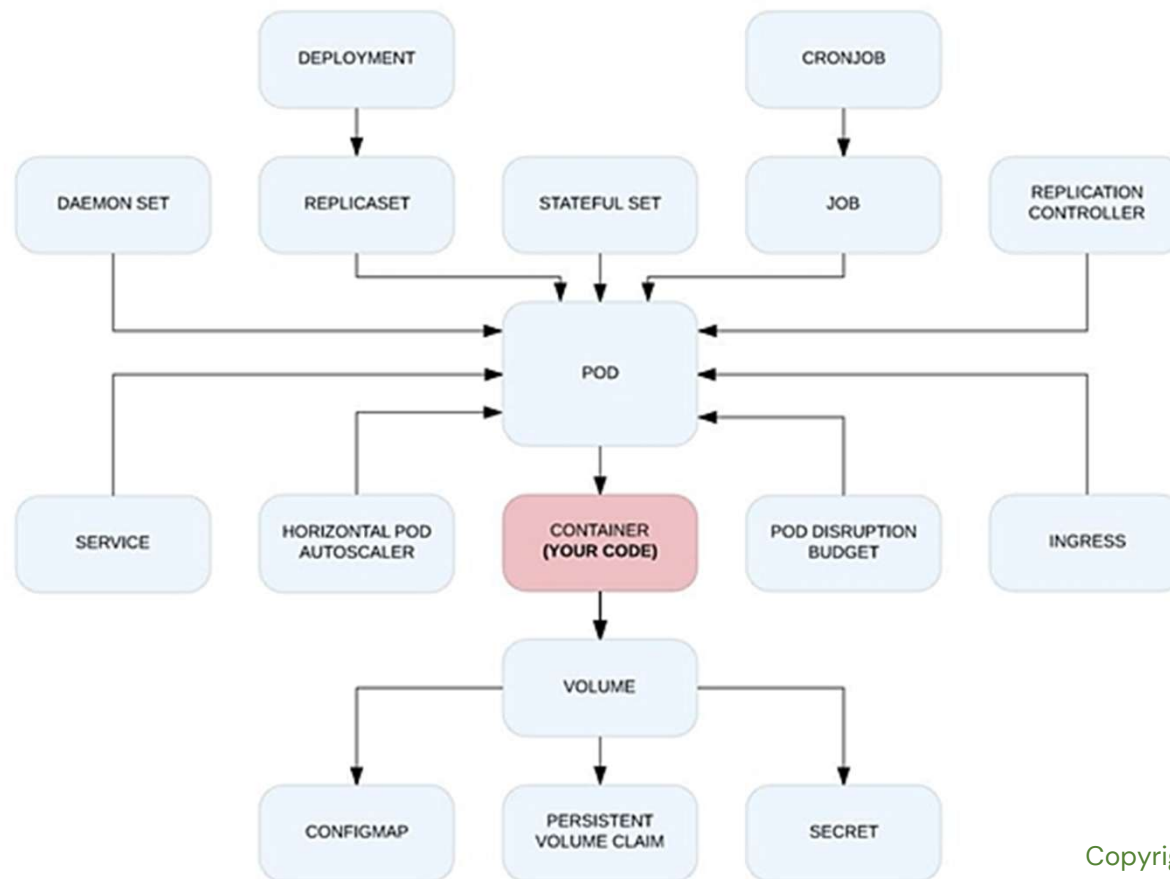
Kubernetes Core Concepts

	Resource (abbr.) [API version]	Description
Security	ServiceAccount (sa) [v1]	An account used by apps running in pods
	Role [rbac.authorization.k8s.io/v1]	Defines which actions a subject may perform on which resources (per namespace)
	ClusterRole* [rbac.authorization.k8s.io/v1]	Like Role, but for cluster-level resources or to grant access to resources across all namespaces
	RoleBinding [rbac.authorization.k8s.io/v1]	Defines who can perform the actions defined in a Role or ClusterRole (within a namespace)
	ClusterRoleBinding* [rbac.authorization.k8s.io/v1]	Like RoleBinding, but across all namespaces

Kubernetes Core Concepts

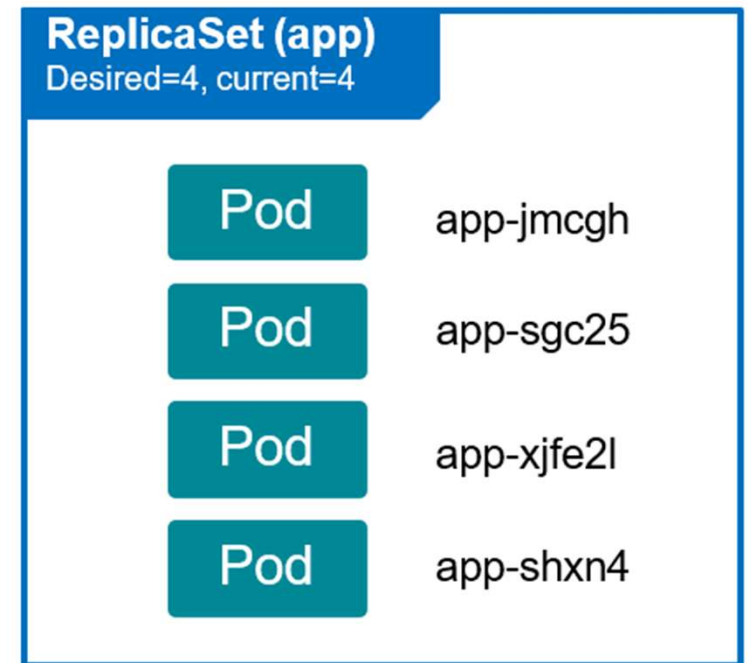
	Resource (abbr.) [API version]	Description
Security	PodSecurityPolicy* (psp) [extensions/v1beta1]	A cluster-level resource that defines which security-sensitive features pods can use
	NetworkPolicy (netpol) [networking.k8s.io/v1]	Isolates the network between pods by specifying which pods can connect to each other
	CertificateSigningRequest* (csr) [certificates.k8s.io/v1beta1]	A request for signing a public key certificate
Ext.	CustomResourceDefinition* (crd) [apiextensions.k8s.io/v1beta1]	Defines a custom resource, allowing users to create instances of the custom resource

Kubernetes Dependencies



Deployment and Release Strategy

- POD Replication – Having many pods with same definition except Name
- Replication Controller – Maintain a stable set of replica pods at any given time



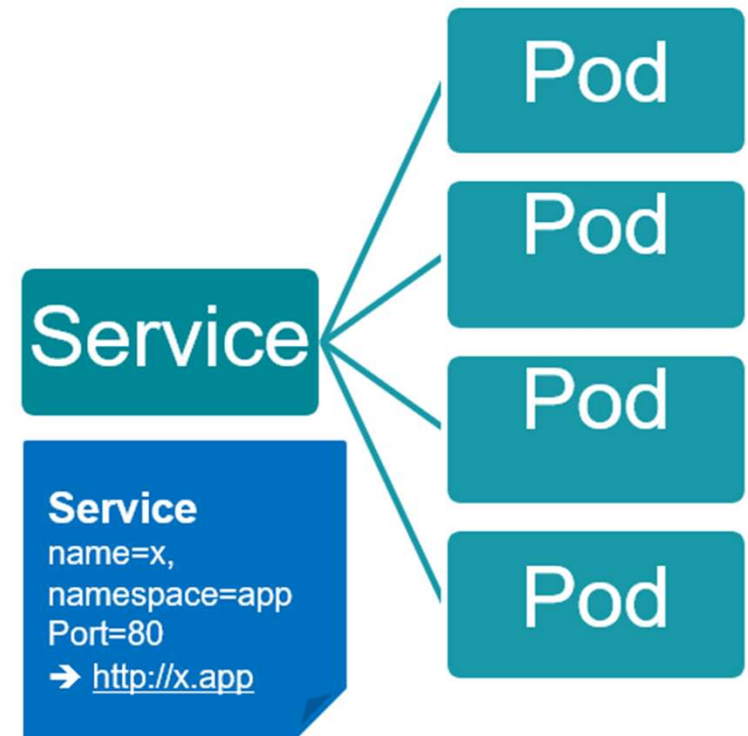
Deployment and Release

- Scaling can be done via
 - Edit YAML file and apply
 - *kubectl edit replicaset*
 - *kubectl scale replicaset ... --replicas X*

Some Background

- All Communication to pods via service
- PODS are ephemeral in nature
- Traffic is routed using labels and selectors

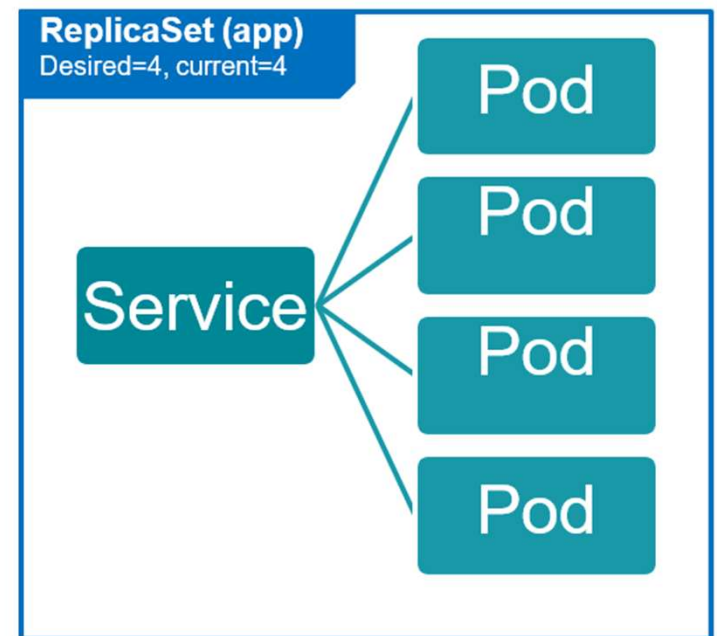
Service Forwards traffic to Pods



Some Background

- ReplicaSet replicates Pods with same Labels
- Because all replicas has same labels, A service with Selectors=Labels can forward traffic in round-robin way.
- Service behaves like loadbalancer for replicas (pods With the same labels)

Service Forwards traffic to



Deployment and Release

- **Scenario:**

- You **deployed** an app few days ago.
- Now you want to **upgrade** your app from v1 to v2.

- **Challenges:**

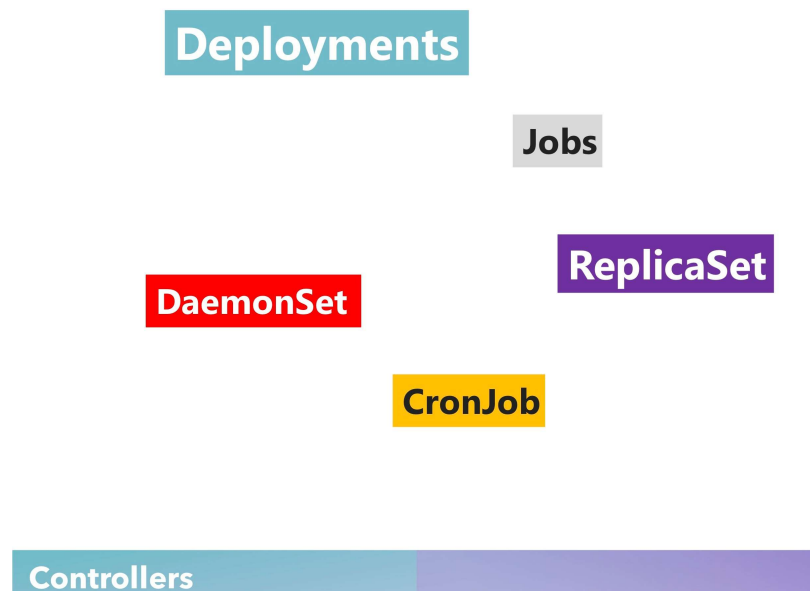
- Can you upgrade with Zero downtime?
- Can you upgrade **replicas** sequentially one after another?
- Can you pause and resume upgrade process?
- Rollback upgrade to previous stable release

Deployment and Release

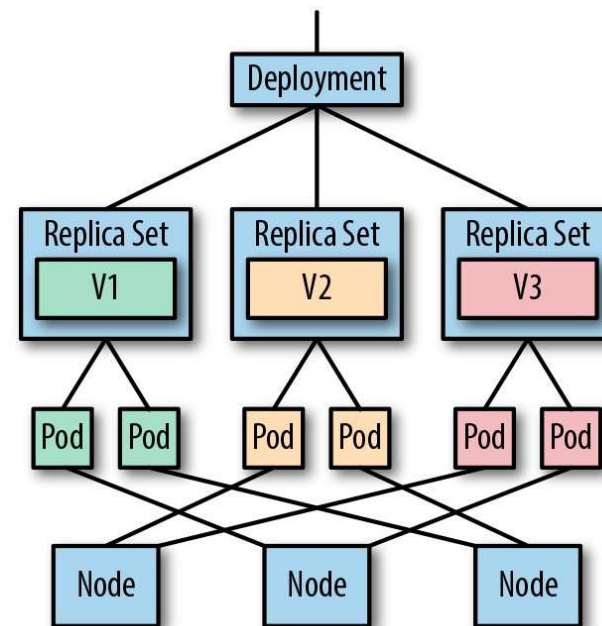
- **In General:**
 - **Deployment** is all of the activities that make a software system available for use
 - Deployment First Release **vs** Deployment Subsequent Releases
 - First Installation **vs** Subsequent Upgrades
- **In Kubernetes:**
 - Deployment is automated using Controllers
 - Deployment controller
 - ReplicaSetController

Deployment and Release

Controllers for Pods



Controller Graphs



Deployment and Release Strategy

- ReplicaSet : Pod Template + Replication
- Deployment: ReplicaSet Template + Update Strategy
- Deployment : Pod Template + Replication + Deployment Strategy
- Name Of Pods Created by Deployment follow this format :
 - *<deployment-name>-<replicaset-name>-<random-string>*

Deployment

- Replication

- Deployment Strategy

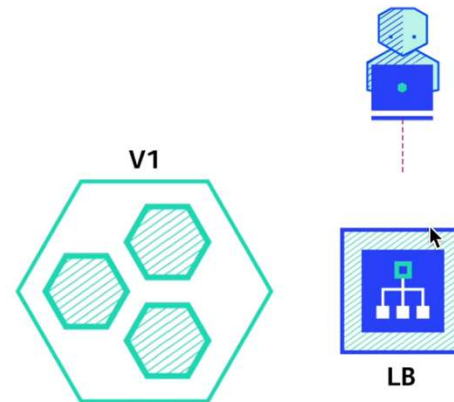
- Pod Template

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: sample
5    labels:
6      section: intro
7  spec:
8    # REPLICATION INFO =====
9    replicas: 3
10   # Deployment Strategy =====
11   # strategy:
12   #   type: RollingUpdate
13   #   rollingUpdate:
14   #     maxUnavailable: 25%
15   #     maxSurge: 25%
16   # Pod Template =====
17   selector:
18     matchLabels: #<- Must match labels of Pod
19     app: sample
20   template: ### PUT HERE POD DEFINITION - except kind and
21     metadata:
22       name: sample
23       labels:
24         app: sample
25     spec:
26       containers:
27       - name: app
28         image: abdenmour/sample-app:v1 # Update v1 to v2
29         ports:
30         - containerPort: 80
```

Deployment and Release Strategy

- **Recreate:**
- Rolling Update (Ramped or incremental)
- Rolling update with Additional
 - batch
- Blue/Green
- Canary
- A/B testing

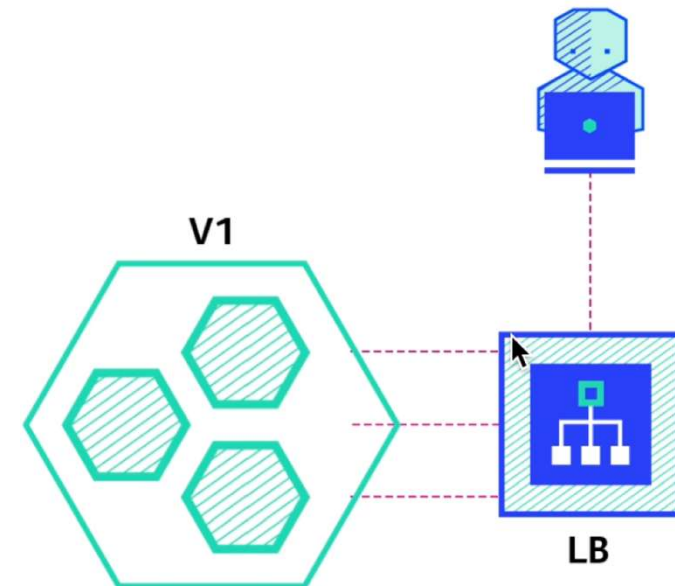
Easy to Setup <> Expect Downtime



Deployment and Release Strategy

- Recreate:
- **Rolling Update** (Ramped or incremental)
- Rolling update with Additional batch
- Blue/Green
- Canary
- A/B testing

Easy to Setup <> Rollout/Rollback take time



Deployment and Release Strategy

- Recreate:
- Rolling Update (Ramped or incremental)
- Rolling update with Additional batch
- Blue/Green
- Canary
- A/B testing

Same as ROLLING UPDATE But :

- Increase number of replicas before start the upgrade

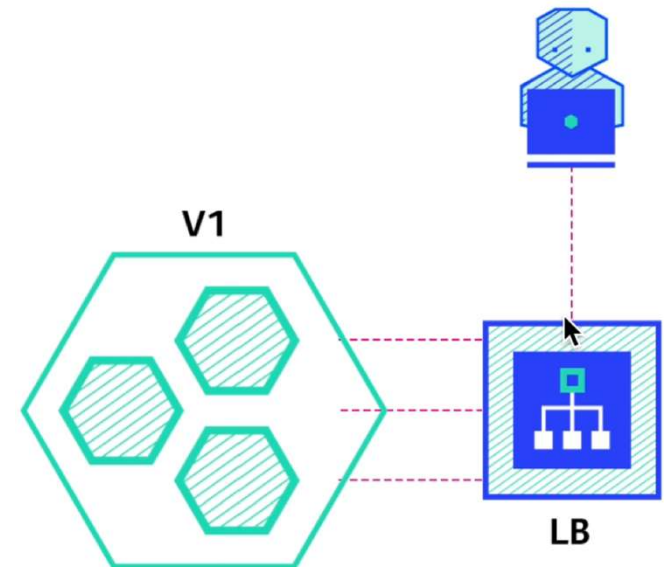
e.g:

- v1 with 3 replicas
- Scale out from 3 to 5 replicas
- Start rolling update v2

Deployment and Release Strategy

- Recreate:
- **Rolling Update** (Ramped or incremental)
- Rolling update with Additional batch
- **Blue/Green**
- Canary
- A/B testing

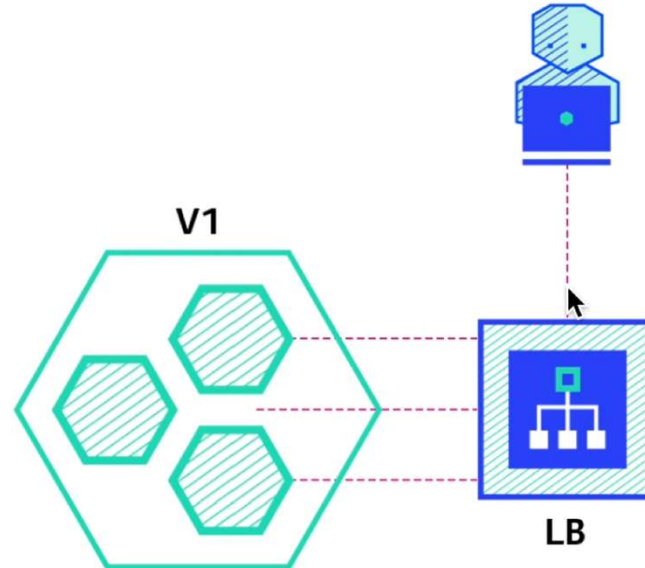
Instant Rollback <> Expensive



Deployment and Release Strategy

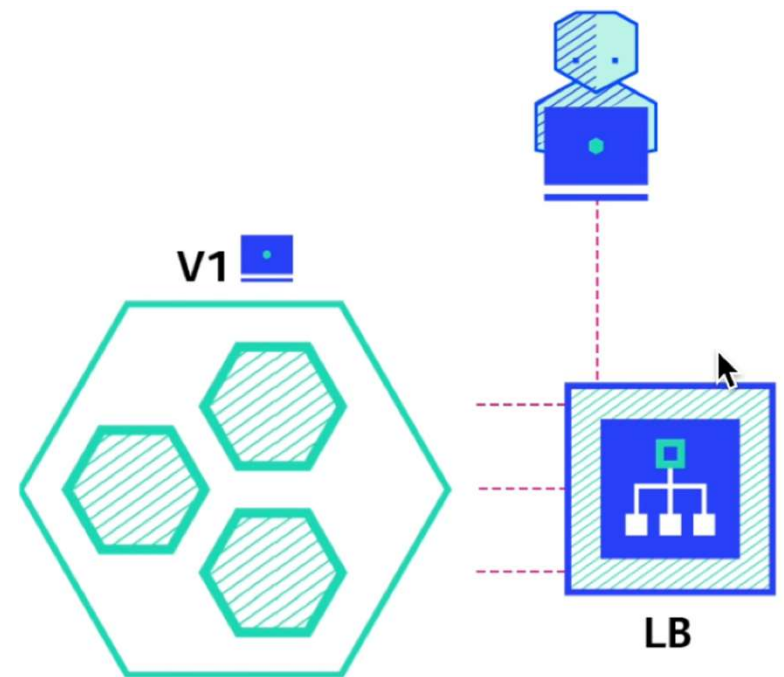
- Recreate:
- **Rolling Update** (Ramped or incremental)
- Rolling update with Additional batch
- **Blue/Green**
- **Canary**
- A/B testing

Fast Rollback <> Slow Rollout



Deployment and Release Strategy

- Recreate:
- **Rolling Update** (Ramped or incremental)
- Rolling update with Additional batch
- Blue/Green
- Canary
- **A/B testing**

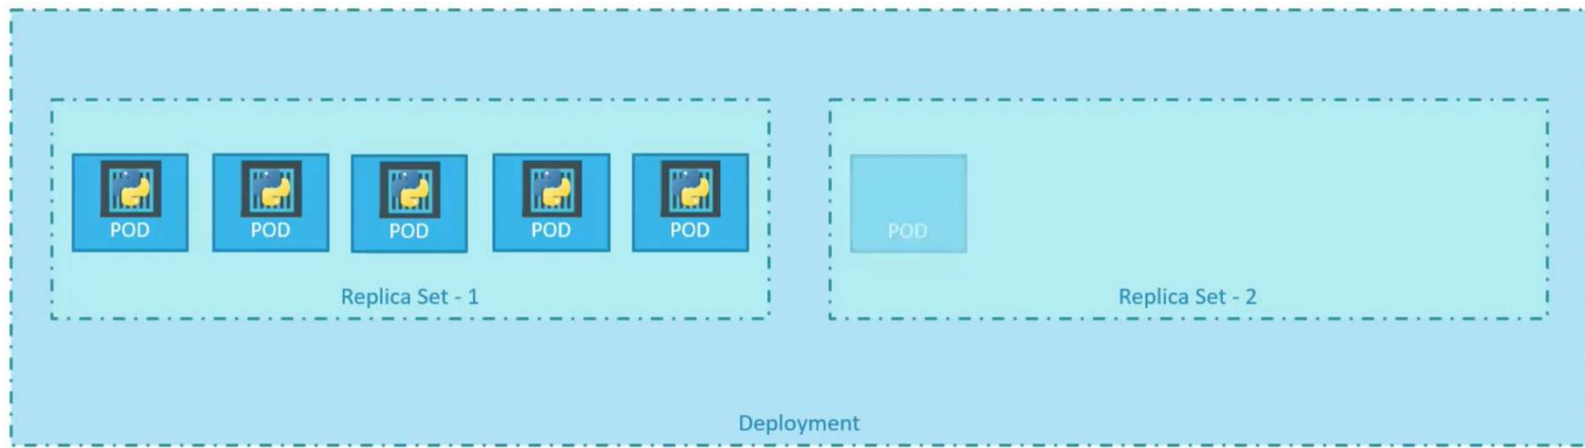


Deploy and Release Strategy

- Recreate
 - Create a new Replica Set V2
 - Scale out Replica Set V2 to max
 - Scale Down Replica set V1 to 0
- Rolling Update or Rolling Update with Additional Batch
 - Gradually Replace existing pods in batches
 - Create new Replica Set V2
 - Scale out Replica set V2 gradually
 - Scale in Replica set v1 gradually
 - Scale in Replica Set V1 to 0

Deploy and Release Strategy

- Rolling Update or Rolling Update with Additional Batch
 - Gradually Replace existing pods in batches
 - Create new Replica Set V2
 - Scale out Replica set V2 gradually
 - Scale in Replica set v1 gradually
 - Scale in Replica Set V1 to 0



A Note – StatefulSets

- A StatefulSet is another Kubernetes controller that manages pods just like Deployments.
- More suited for stateful apps.
- A stateful application requires pods with a unique identity (for example, hostname). One pod should be able to reach other pods with well-defined names.
- For a StatefulSet to work, it needs a Headless Service. A Headless Service does not have an IP address.
- Internally, it creates the necessary endpoints to expose pods with DNS names. The StatefulSet definition includes a reference to the Headless Service, but you have to create it separately.

A Note – StatefulSets

- By nature, a StatefulSet needs persistent storage so that the hosted application saves its state and data across restarts. Kubernetes provides Storage Classes, Persistent Volumes, and Persistent Volume Claims to provide an abstraction layer above the cloud provider's storage-offering mechanism.
- Once the StatefulSet and the Headless Service are created, a pod can access another one by name prefixed with the service name.

Design Patterns – Single Container

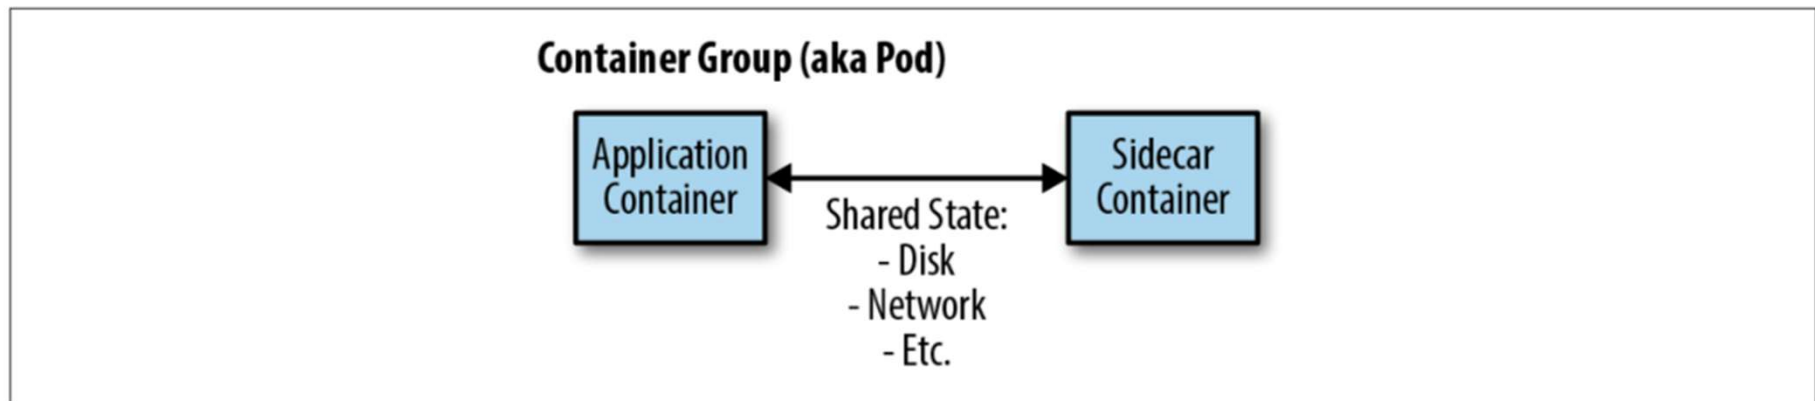
Single node base for other patterns

- Establish (reserve) resource boundaries
- Isolate resources from others (protect)
- Separation of concerns (single responsibility)

Design Patterns – Side Car

Application Container

Extend the application (SideCar)

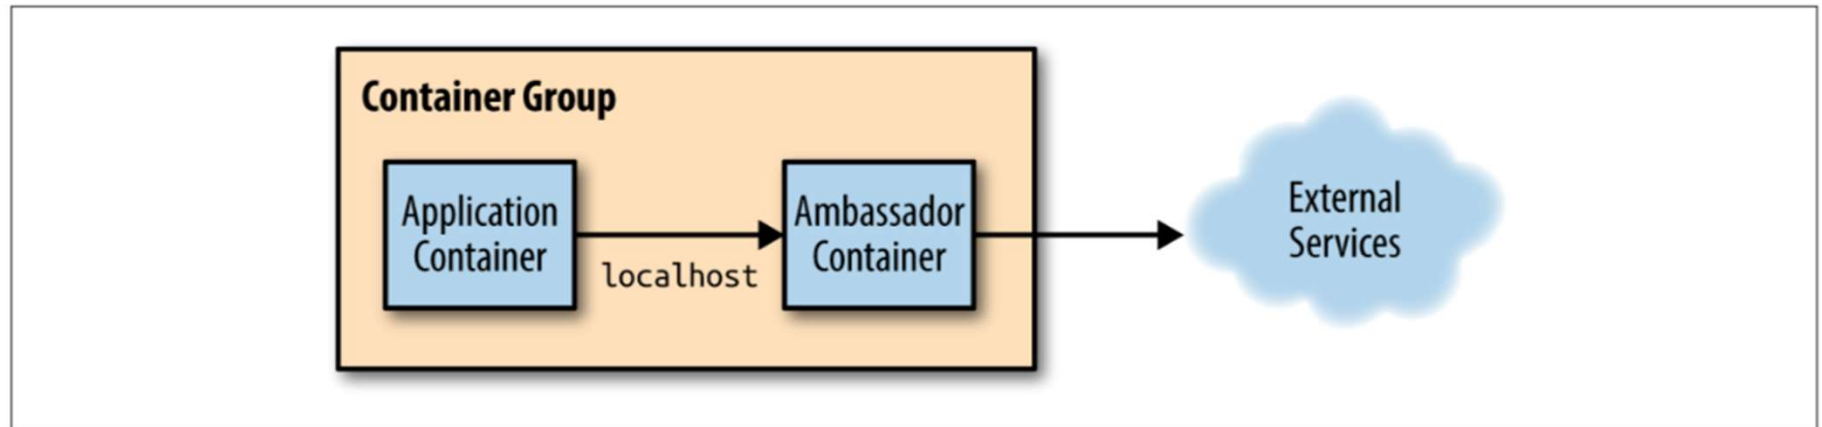


Design Patterns – Side Car

- Co-scheduled in the same node
- Better performance when communicating
- Share resources (volumes)
- Examples:
 - HTTPS Layer
 - HTTP Proxy for Self-hosted apps

Design Patterns – Ambassador

- Application container
- **Broker** between application and consumers
- Grouped as with sidecar pattern

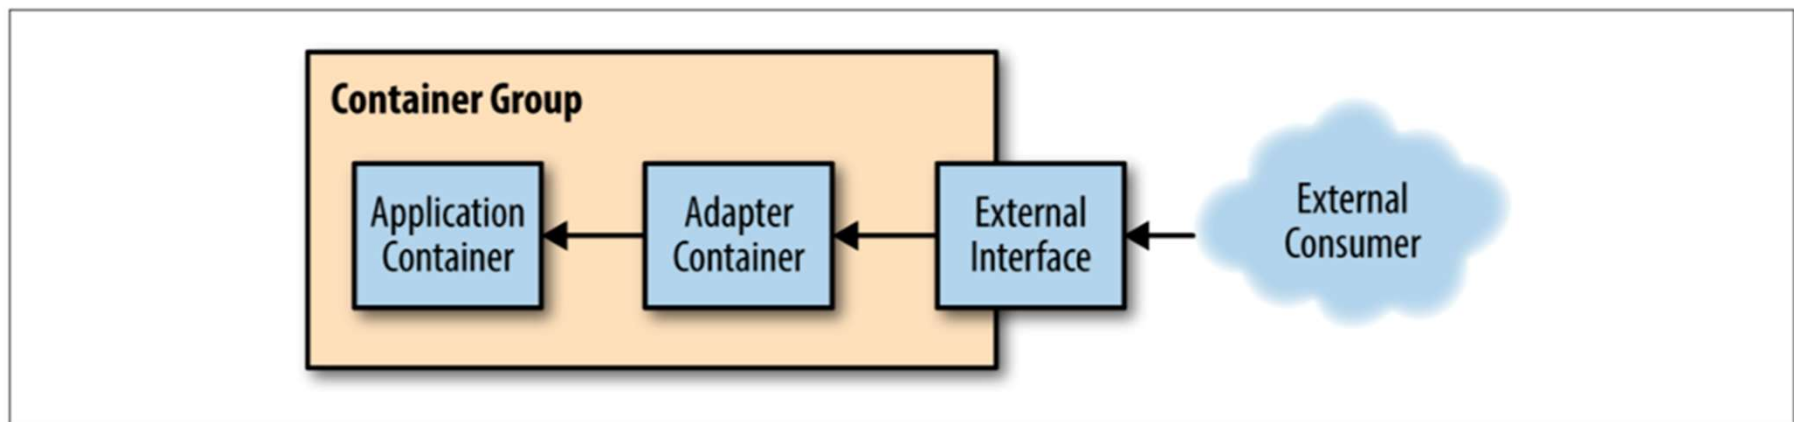


Design Patterns – Ambassador

- Modular and reusable containers
- Reuse in other applications and environments
- Like a load balancer (tradeoffs)
- Examples:
 - Shard a service (service discovery)
 - Service brokering (IaaS or PaaS DBs)
 - A/B Testing (split traffic)

Design Patterns – Adapter

- Application container
- Modify the interface of the application



Design Patterns – Adapter

- Helps to set a standard in communication
- Adapts to provide a consistent interface
- Reusable modules (open source, internal, etc.)
- Examples:
 - Collect metrics
 - Parse and store logs (like Logstash, Fluentd)
 - Custom health checks of the application

Design Patterns – Overview

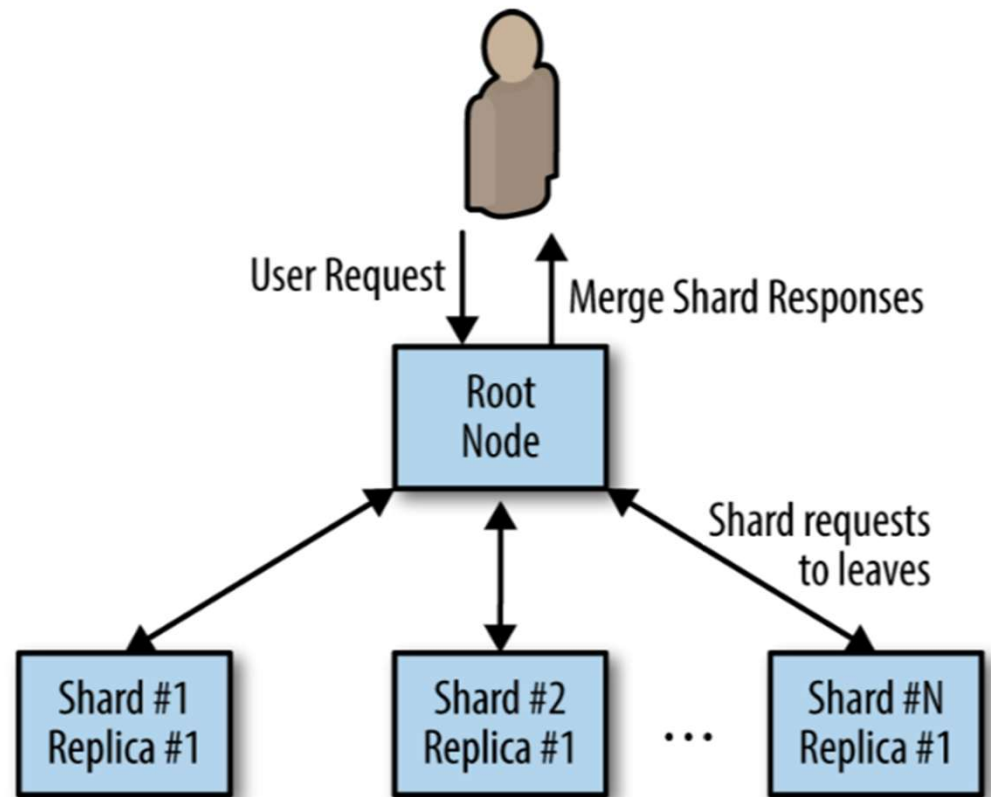
- Single Container: Isolate, Portable, etc.
- Sidecar: Extend behavior.
- Ambassador: Broker or proxy to the exterior.
- Adapter: Consistent communication.

Design Patterns – Multi-container

- Scale at the Node Level
- Loosely Coupled
- Communication with Network Calls
- Parallel Calls Natively
- Base for Microservices (stateless)

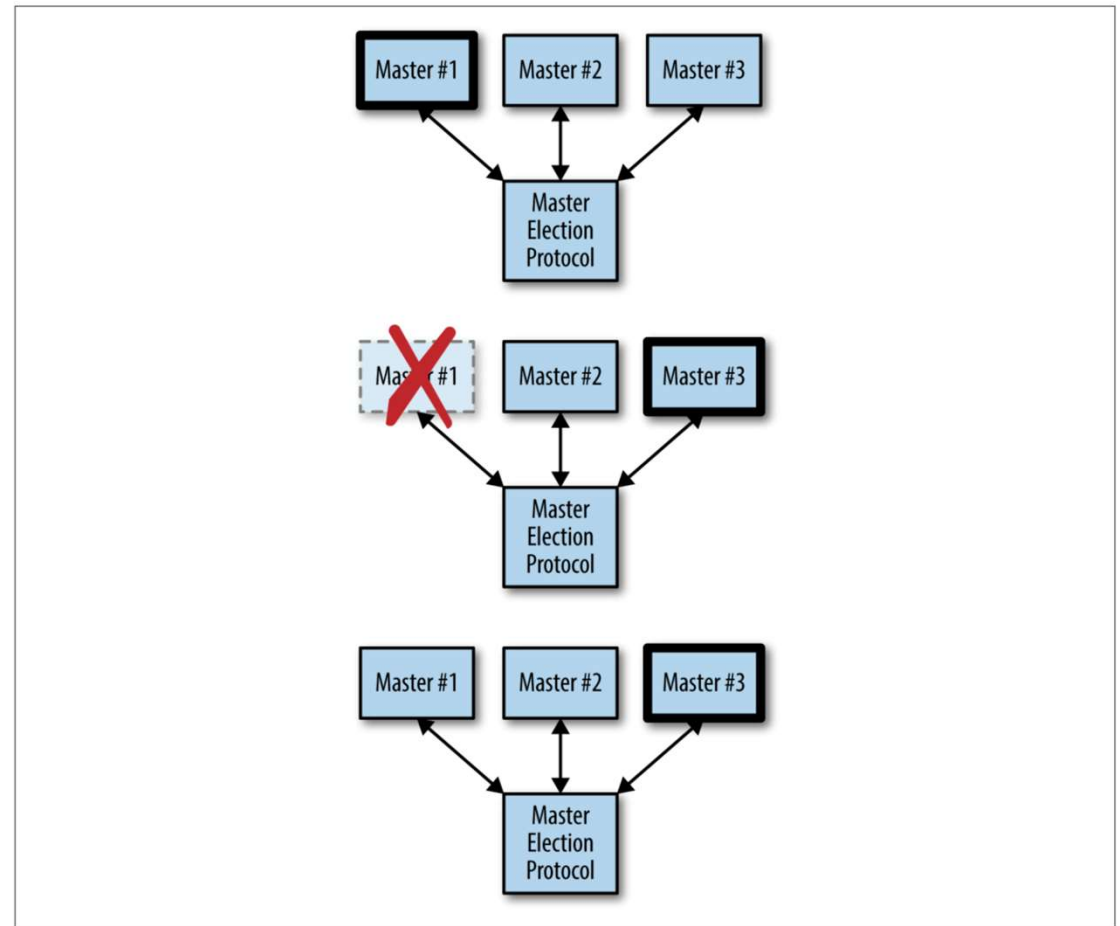
Design Patterns – Scatter/Gather

- Split a big task
- Simultaneously
- Combined results



Design Patterns – Leader/Election

- Master #1 is selected
- Master #1 fails
- Master #3 is selected

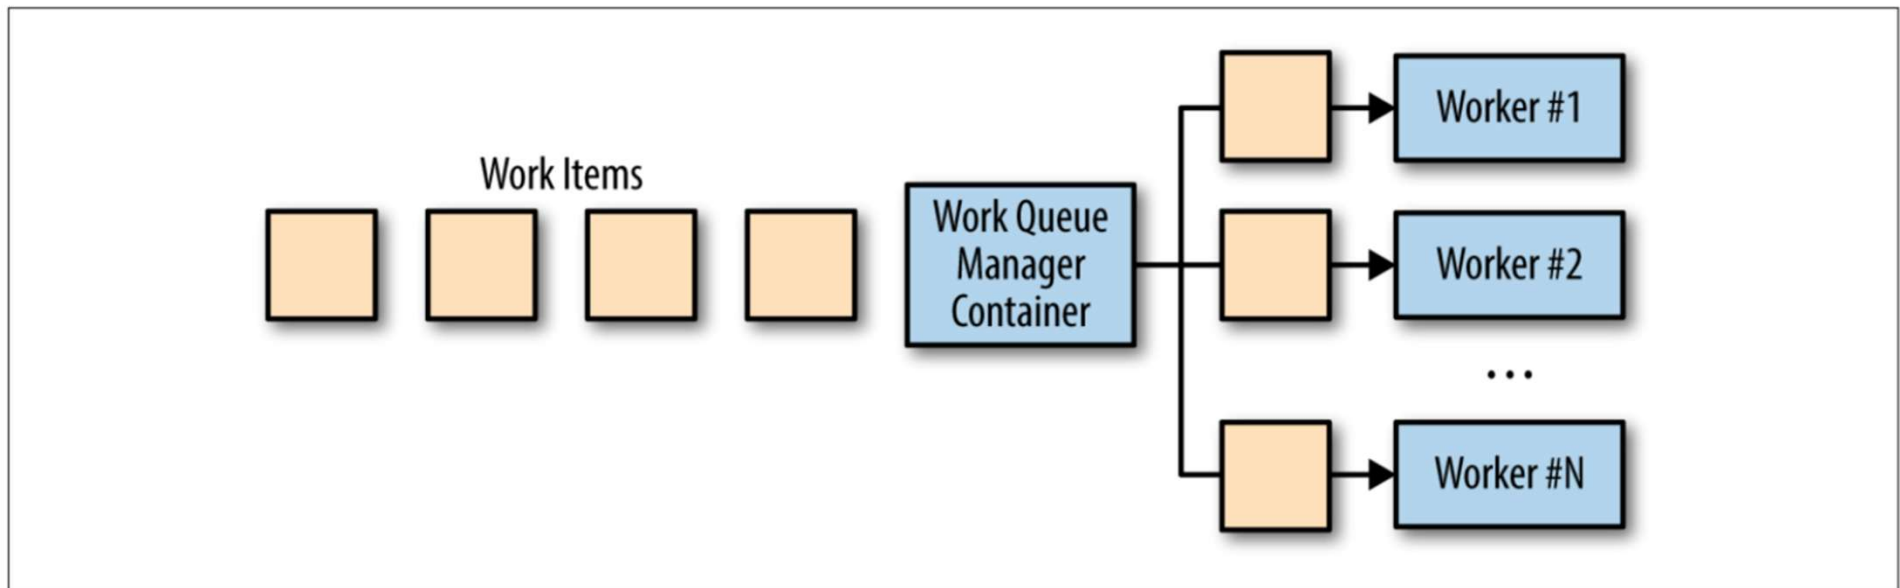


Design Patterns – Leader/Election

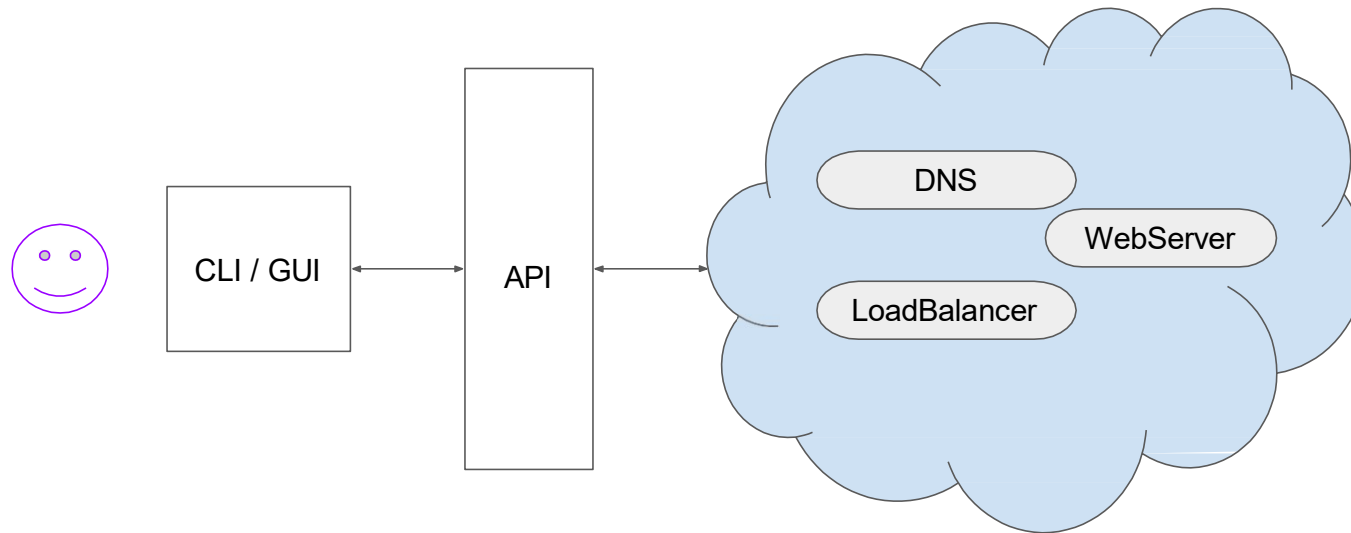
- With Kubernetes:
 - If container crash, it's recreated
 - If container hangs, it's restarted (health check)
 - If node fails, container it's re-scheduled
- But applications might need a high SLA
- Reuse, don't reinvent the wheel
- Algorithms like Paxos or RAFT
- Container implementations
 - etcd
 - Zookeeper
 - consul

Design Patterns – Work Queue

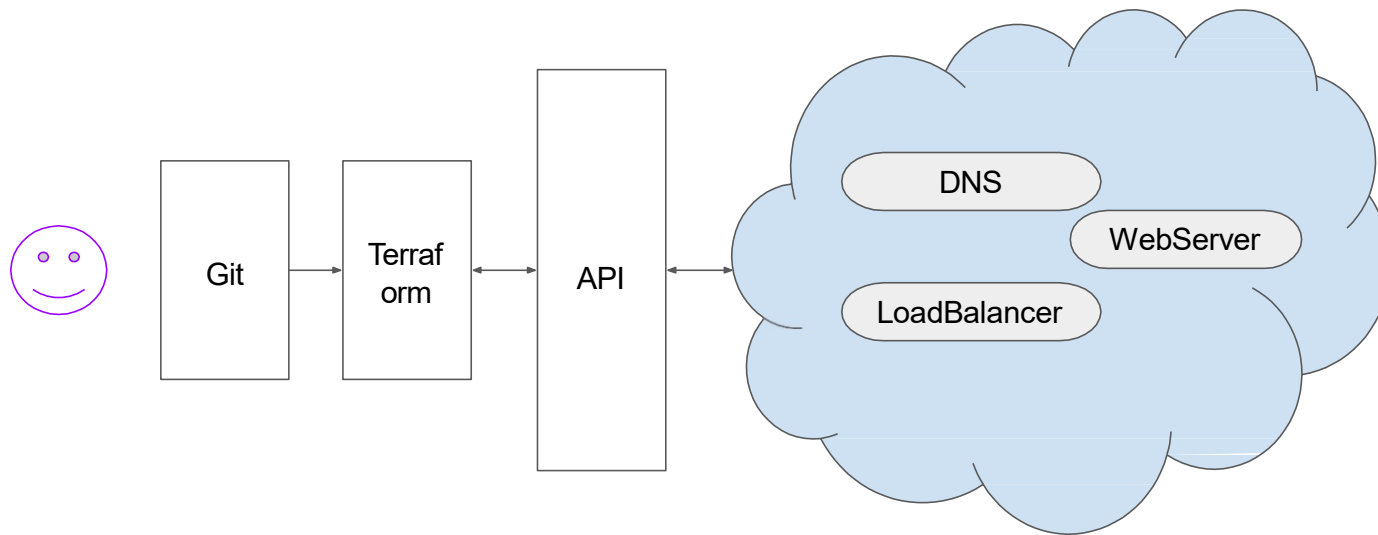
- Manager Container (one)
- Worker Container (multiple)



Infra Provisioning – Manual

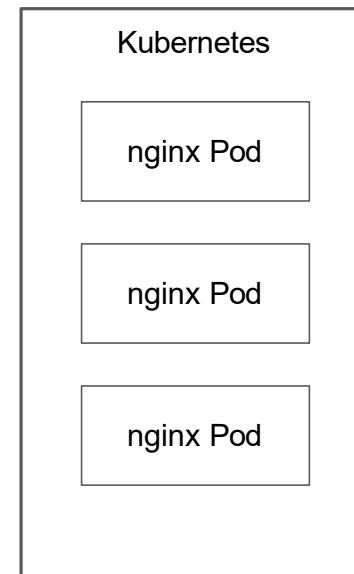
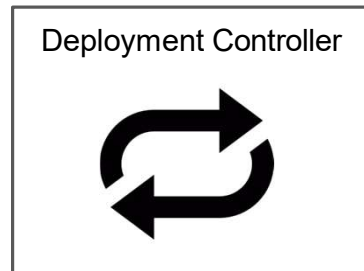


Infra Provisioning – Configuration Mgmt.



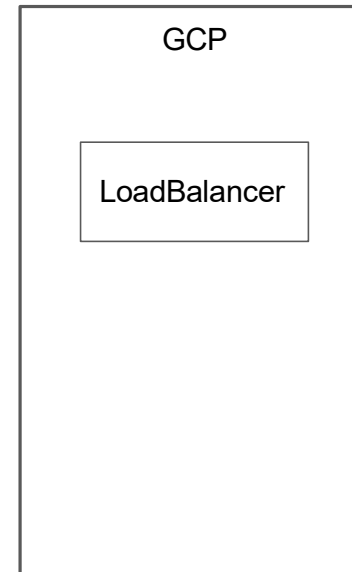
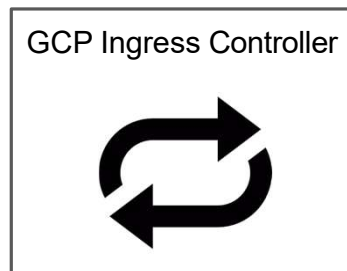
Kubernetes Controllers

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```



Kubernetes Controllers

```
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: test-ingress
spec:
  rules:
  - http:
      paths:
      - path: /testpath
        backend:
          serviceName: test
          servicePort: 80
```



Kubernetes Custom Resources

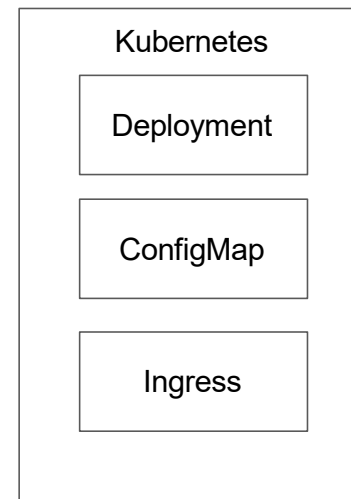
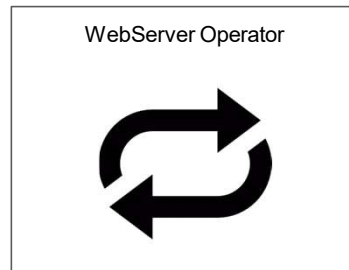
- Extensions of the Kubernetes API
- Define your own “Kind”
- Use capabilities of kubernetes objects. (RBAC, Finalizers, Versioning, Watches,...)
- Kubectl support

Kubernetes Custom Resources

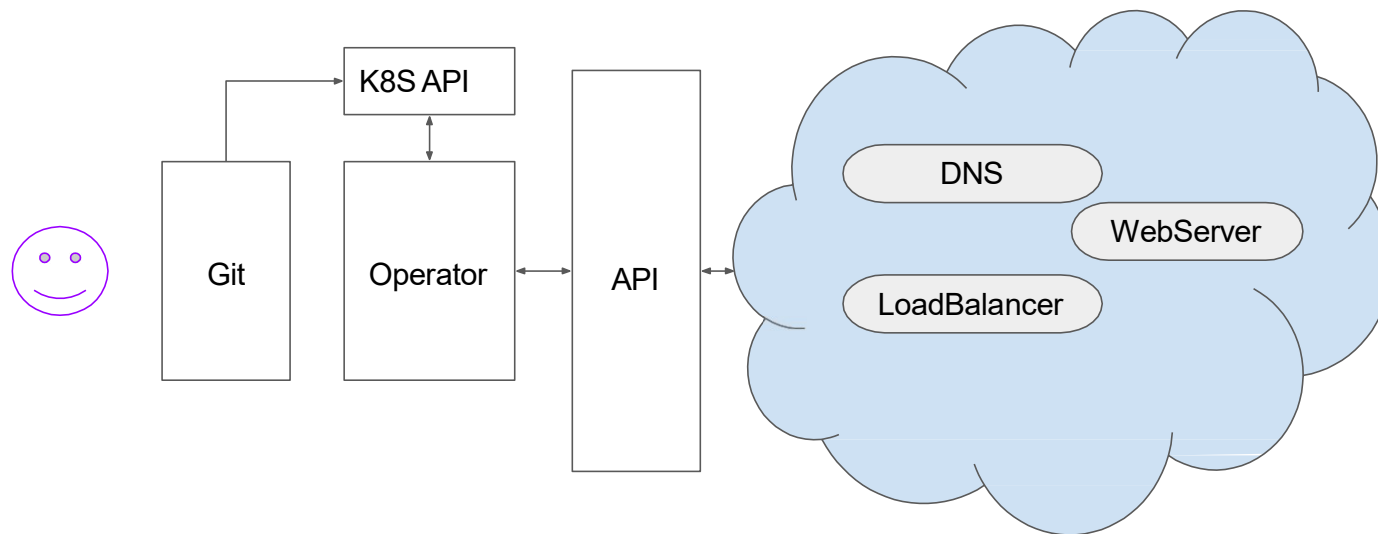
```
apiVersion: webser.attila.com/v1
kind: WebServer
metadata:
  name: mywebserver
spec:
  www: <html>hello world</html>
status:
  url: http://34.2.44.5/myweb
```

Kubernetes Custom Resources

```
apiVersion: webser.attila.com/v1
kind: WebServer
metadata:
  name: mywebserver
spec:
  www: <html>hello world</html>
status:
  url: http://34.2.44.5/myweb
```



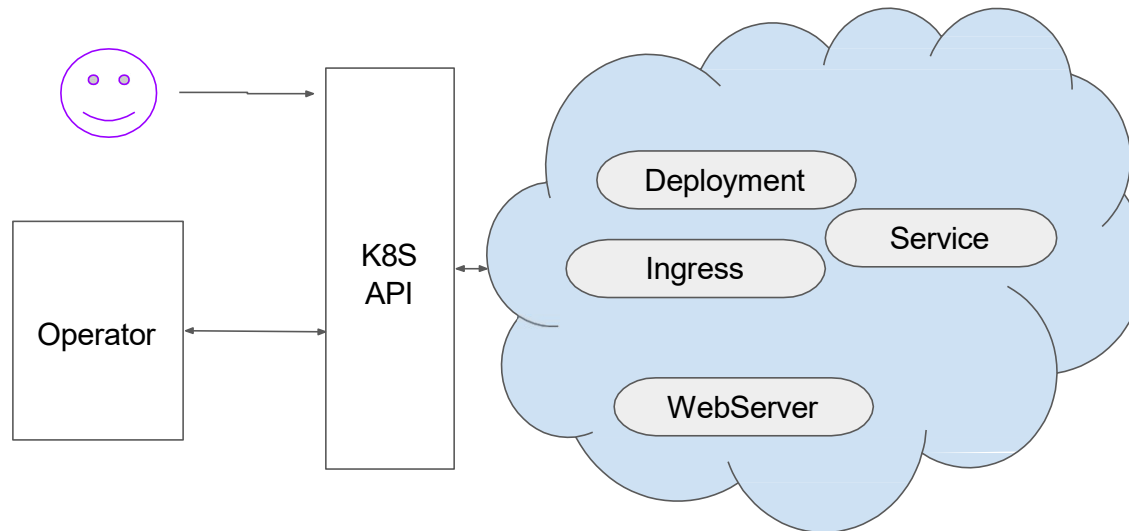
Infra Provisioning – Operators



Operators

- Just a pattern - Kubernetes doesn't know about your Operator
- Just a pattern - not trivial to implement
- It's how Kubernetes Controllers work
- We have frameworks: operator-sdk (Go), java-operator-sdk (JDK)

Infra Provisioning – Operators



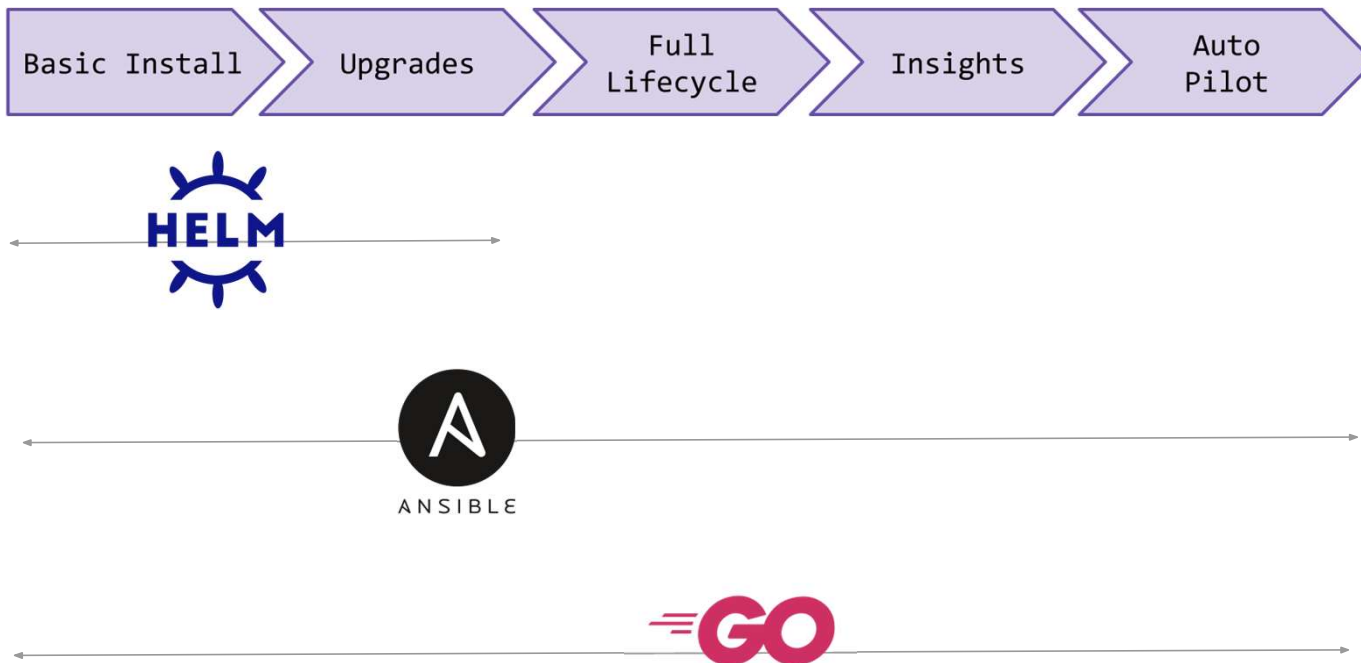
Operators

- Operators are a very robust and flexible way to write provisioning logic
- Use case 1: Provision stateful resources on Kubernetes
- Use-case 2: Provision complex systems

Operators

- Deploying an application on Demand
- Taking and Restoring Backups of app State
- Handling Upgrades including DB Schema changes and/or config setting
- Publish Services to Applications that do not support Kubernetes API to discover them
- Simulating Failure in all or parts of your system
- Choosing Leader for distributed apps without election process

Operators



Deploying Apps

Deploy Applications with and without Helm



We're done!
Thank you for your time and
participation.